

Linear Regression

Data Mining

Prof. Sujee Lee

Department of Systems Management Engineering


Sungkyunkwan University

Linear Regression

Linear Regression

■ Linear Regression

- Linear regression models $\hat{y} = f(\mathbf{x})$ as a linear function of input features.



Learning a Linear Regression Model

- **Linear regression on single variate data** Assume $y_i \in \mathbb{R}, x_i \in \mathbb{R}$

- Goal : Given data pairs $\{(x_i, y_i)\}_{i=1}^n$, find a “best fit” line through data

- Model : $\hat{y} = f(x) = ax + b$

- Least Squares :

$$\min_{a,b} \sum_{i=1}^n (y_i - (ax_i + b))^2$$

Learning a Linear Regression Model

- To solve the following minimization problem (optimization problem)

$$\min_{a,b} \sum_{i=1}^n (y_i - (ax_i + b))^2$$

- take derivatives :
- System of linear equations :
- Solution :

$$a^* = \frac{n \sum_{i=1}^n x_i y_i - (\sum_{i=1}^n x_i) (\sum_{i=1}^n y_i)}{n \sum_{i=1}^n x_i^2 - (\sum_{i=1}^n x_i)^2}, \quad b^* = \frac{\sum_{i=1}^n y_i - a \sum_{i=1}^n x_i}{n}$$

Learning a Linear Regression Model

- Q: How do we generalize to more dimensions?
- **Linear regression on multi variate data**
 - $y_i \in \mathbb{R}, \mathbf{x}_i = (x_{i,1}, x_{i,1}, \dots, x_{i,d}) \in \mathbb{R}^d$
 - Goal : find best linear predictor of y_i using all components of vector \mathbf{x}_i
 - use $w_1 x_{i,1} + \dots + w_d x_{i,d} + b$ to predict y_i
- we need to solve

$$\min_{w_0, w_1, w_2, \dots, w_d} \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

Let $f(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2$

Take a gradient of $f(\mathbf{w})$ and set it equals to 0 : $\nabla_{\mathbf{w}} f(\mathbf{w}) = 0$

$$\begin{aligned} \nabla_{\mathbf{w}} \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2 &= \sum_{i=1}^n \nabla_{\mathbf{w}} (y_i - \mathbf{w}^T \mathbf{x}_i)^2 \\ &= \sum_{i=1}^n 2(y_i - \mathbf{w}^T \mathbf{x}_i) \nabla_{\mathbf{w}} (y_i - \mathbf{w}^T \mathbf{x}_i) \\ &= \sum_{i=1}^n -2(y_i - \mathbf{w}^T \mathbf{x}_i) \nabla_{\mathbf{w}} \mathbf{w}^T \mathbf{x}_i \end{aligned}$$

gradient of $\mathbf{w}^T \mathbf{x}_i$, compute partial derivative w.r.t each w_j

$$\frac{\partial}{\partial w_j} (\mathbf{w}^T \mathbf{x}_i) = \frac{\partial}{\partial w_j} \left(\sum_{k=0}^d w_k x_{i,k} \right) = x_{i,j}, \forall j$$

$$\nabla_{\mathbf{w}} \mathbf{w}^T \mathbf{x}_i = \mathbf{x}_i$$

Thus, $\nabla_{\mathbf{w}} f(\mathbf{w}) = 0 \Leftrightarrow \sum_{i=1}^n -2(y_i - \mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i = 0$

Solve for \mathbf{w} ?

$$\sum_i -2(y_i - \mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i = 0$$

$$\sum_i (\mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i = \sum_i y_i \mathbf{x}_i$$

$$(\sum_i \mathbf{x}_i \mathbf{x}_i^T) \mathbf{w} = \sum_i y_i \mathbf{x}_i$$

Ordinary Least Square estimate:

$$\therefore \mathbf{w}_{OLS} = \left(\sum_i \mathbf{x}_i \mathbf{x}_i^T \right)^{-1} \cdot \sum_i y_i \mathbf{x}_i$$

Does it agree with 'a' and 'b' in 1-dim case?

$$\hat{y} = \mathbf{w}_{OLS}^T \mathbf{x}$$

Learning a Linear Regression Model

■ Matrix Representation

Consider the matrix & vector $\mathbf{X} \in \mathbb{R}^{n \times (d+1)}, \mathbf{y} \in \mathbb{R}^n$

Least squares objective function can be written as

$$\|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2$$

Recall: if $\mathbf{v} \in \mathbb{R}^n$, then $\|\mathbf{v}\|_2^2 = \sum_{j=1}^n v_j^2$

Now, we want to $\min_{\mathbf{w}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2$

want to take a gradient of f :

$$\begin{aligned}\nabla_{\mathbf{w}} f(\mathbf{w}) &= \nabla_{\mathbf{w}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 = \nabla_{\mathbf{w}} (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) \\ &= \nabla_{\mathbf{w}} [\mathbf{y}^T \mathbf{y} - (\mathbf{X}\mathbf{w})^T \mathbf{y} - \mathbf{y}^T \mathbf{X}\mathbf{w} + (\mathbf{X}\mathbf{w})^T (\mathbf{X}\mathbf{w})] \\ &= \nabla_{\mathbf{w}} [\mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T \mathbf{X}\mathbf{w} + \mathbf{w}^T \mathbf{X}^T \mathbf{X}\mathbf{w}] \\ &= -2\mathbf{X}^T \mathbf{y} + \nabla_{\mathbf{w}} (\mathbf{w}^T \mathbf{X}^T \mathbf{X}\mathbf{w})\end{aligned}$$

Claim: For any symmetric matrix \mathbf{A} , $\nabla_{\boldsymbol{\beta}} (\boldsymbol{\beta}^T \mathbf{A} \boldsymbol{\beta}) = 2\mathbf{A}\boldsymbol{\beta}$

$$\nabla_{\mathbf{w}} f(\mathbf{w}) = -2\mathbf{X}^T \mathbf{y} + 2\mathbf{X}^T \mathbf{X}\mathbf{w} = 0$$

$$\nabla_{\mathbf{w}} f(\mathbf{w}) = -2\mathbf{X}^T \mathbf{y} + 2\mathbf{X}^T \mathbf{X} \mathbf{w} = 0$$

Solve for \mathbf{w} :

$$\mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$\therefore \mathbf{w}_{OLS} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \qquad \mathbf{w}_{OLS} = \left(\sum_i \mathbf{x}_i \mathbf{x}_i^T \right)^{-1} \cdot \sum_i y_i \mathbf{x}_i$$

Linear Regression

- **Linear Regression (ordinary least squares (OLS))**

- Linear regression finds the parameters \mathbf{w} and b that minimize the mean squared error between predictions and the true regression targets on the training set.

$$\hat{y} = \mathbf{w}^T \mathbf{x} + b = w_1 x_1 + \cdots + w_d x_d + b$$

$$\mathbf{x} = (x_1, \dots, x_d) \in \mathbb{R}^d, \quad y, \hat{y} \in \mathbb{R}$$

- It solves

$$\min_{\mathbf{w}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2$$

- The solution is

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- The trained model is

$$f(\mathbf{x}) = \mathbf{w}^{*T} \mathbf{x}$$

- Linear regression has no hyperparameters, thus has no way to control model complexity.

Other Views

Linear Regression

- **Linear Regression (ordinary least squares (OLS))**

- Linear regression solves

$$\min_w \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2$$

- The solution is

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- The trained model is

$$f(\mathbf{x}) = \mathbf{w}^{*T} \mathbf{x}$$

- **Three views of ordinary least squares**

- Algebraic (matrices, gradients)
- Geometric
- Probabilistic

Probabilistic Approach for OLS

■ Probabilistic Approach for OLS

- Idea: Instead of thinking just in terms of data points $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$ or data matrices (\mathbf{X}, \mathbf{y}) , considers a **generative model** (probabilistic)
- Consider a setting where (\mathbf{x}_i, y_i) are generated in a random way:

$$y_i = \mathbf{w}^T \mathbf{x}_i + \epsilon_i, \text{ where } \mathbf{x}_i \text{'s are fixed vectors and } \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

ϵ_i is normally distributed in \mathbb{R} , with variance σ^2

- Start with \mathbf{x}_i , generate random ϵ_i , which gives y_i
- Goal : Estimate \mathbf{W} using the observations $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$

Probabilistic Approach for OLS

■ Maximum Likelihood Estimation (MLE)

- We will use maximum likelihood estimation (MLE) to find our estimate of \mathbf{w}
- We write down the probability of seeing $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$, assuming \mathbf{w} was true regression vector, and maximize it over \mathbf{w}

Likelihood (probability of seeing the data):

$$L_{\mathbf{w}}(\mathbf{X}, \mathbf{y}) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - \mathbf{w}^T \mathbf{x}_i)^2}{2\sigma^2}\right)$$

Recall: p.d.f. of a random variable x following $\mathcal{N}(\mu, \sigma^2)$ is

$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

$$\begin{aligned} L_{\mathbf{w}}(\mathbf{X}, \mathbf{y}) &= P(\{(\mathbf{x}_i, y_i)\}_{i=1}^n) \\ &= \prod_i^n P(\mathbf{x}_i, y_i) \\ &= \prod_i^n P(\epsilon_i = y_i - \mathbf{w}^T \mathbf{x}_i) \\ &= \prod_i^n p(y_i - \mathbf{w}^T \mathbf{x}_i) \\ &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - \mathbf{w}^T \mathbf{x}_i)^2}{2\sigma^2}\right) \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma}\right)^n \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2\right) \end{aligned}$$

Now, we want to maximize likelihood $L_{\mathbf{w}}(\mathbf{X}, \mathbf{y})$ with respect to \mathbf{w}

$$L_{\mathbf{w}}(\mathbf{X}, \mathbf{y}) = \left(\frac{1}{\sqrt{2\pi}\sigma} \right)^n \exp \left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2 \right)$$

Maximizing likelihood $L_{\mathbf{w}}(\mathbf{X}, \mathbf{y})$ with respect to \mathbf{w} is the same as minimizing $\sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2$ with respect to \mathbf{w} (**this is the same as OLS**)

Pros:

- We get a whole family of estimators (for different distributions of)
- We can do inference (i.e, confidence intervals, hypothesis, tests, etc.)

Cons:

- Assumptions on how data are generated (linear relationship, Gaussian errors)