

Tree Ensemble Models

Data Mining

Prof. Sujee Lee

Department of Systems Management Engineering

Sungkyunkwan University

Decision Tree

■ The main hyperparameters of decision trees

- Picking one of the pre-pruning strategies (`max_depth`, `max_leaf_nodes`, or `min_samples_leaf`) is sufficient to prevent overfitting.
- * Typically chosen to have the highest performance in validation data.

■ Strengths

- Decision trees work well when you have a mix of continuous and categorical features.
- The algorithms are completely invariant to scaling of the data. (no data scaling is needed)
- Feature selection & reduction is automatic.
- It is robust to noise.
- The resulting model can easily be visualized and understood.

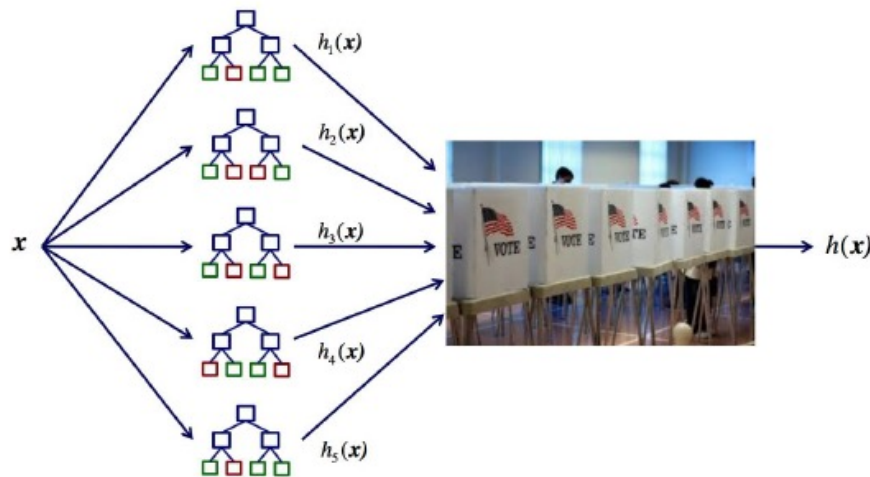
■ Weaknesses

- Even with the use of pre-pruning, they tend to overfit and provide poor generalization performance.
 - Thus, the ensemble methods are usually used in place of a single decision tree.

Tree Ensembles

■ Tree Ensembles

- Classification trees get overfitted easily (low bias – high variance) and are seldom accurate
- Thus, we will build multiple trees and average their results!
- **Ensemble** is a way of averaging multiple deep decision trees, trained on different parts of the same training set, to reduce the variance
 - combine the outputs of many “weak” classifiers to produce a powerful “committee”



regression :

$$f(x) = \frac{1}{B} \sum_{b=1}^B f_b(x)$$

classification :


$$C(x) = \text{majority vote } \{C_b(x)\}_1^B$$

Bagging, Boosting, Random Forest

Bagging

■ Bagging (Bootstrap Aggregating)


- Each tree might do a relatively good job of predicting, but will overfit on part of the data in different ways.
- If we build many trees, we can reduce the amount of overfitting (reduce the variance of the models) by averaging their results while retaining the predictive power of the trees.
- In Bagging, it simply
 - repeat [choose different subsamples (bootstrap samples) + fit a model]
 - take the majority vote
- Not only for trees!
 - works well for high-variance, low-bias models



Boosting

■ Boosting


- Unlike Bagging, the committee of weak learners evolves over time (sequentially learning trees), and the members cast a weighted vote.
- Boosting appears to dominate bagging on most problems and is preferred.



Boosting

■ AdaBoost (Adaptive Boosting)

- Combine multiple weak learners (decision trees with a single split)
- In each step, give more weights on observations misclassified in the previous step




Algorithm 10.1 *AdaBoost.M1*

1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.
 2. For $m = 1$ to M :
 - (a) Fit a classifier $G_m(x)$ to the training data using weights w_i .
 - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
 - (c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.
 - (d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.
 3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.
-

Boosting

■ Gradient Boosting

- Try to fit the new classifier to the residual errors made by the previous classifier



Algorithm 10.3 Gradient Tree Boosting Algorithm.

1. Initialize $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$.

2. For $m = 1$ to M :

(a) For $i = 1, 2, \dots, N$ compute

$$r_{im} = - \left[\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}.$$

(b) Fit a regression tree to the targets r_{im} giving terminal regions R_{jm} , $j = 1, 2, \dots, J_m$.

(c) For $j = 1, 2, \dots, J_m$ compute

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma).$$

(d) Update $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$.

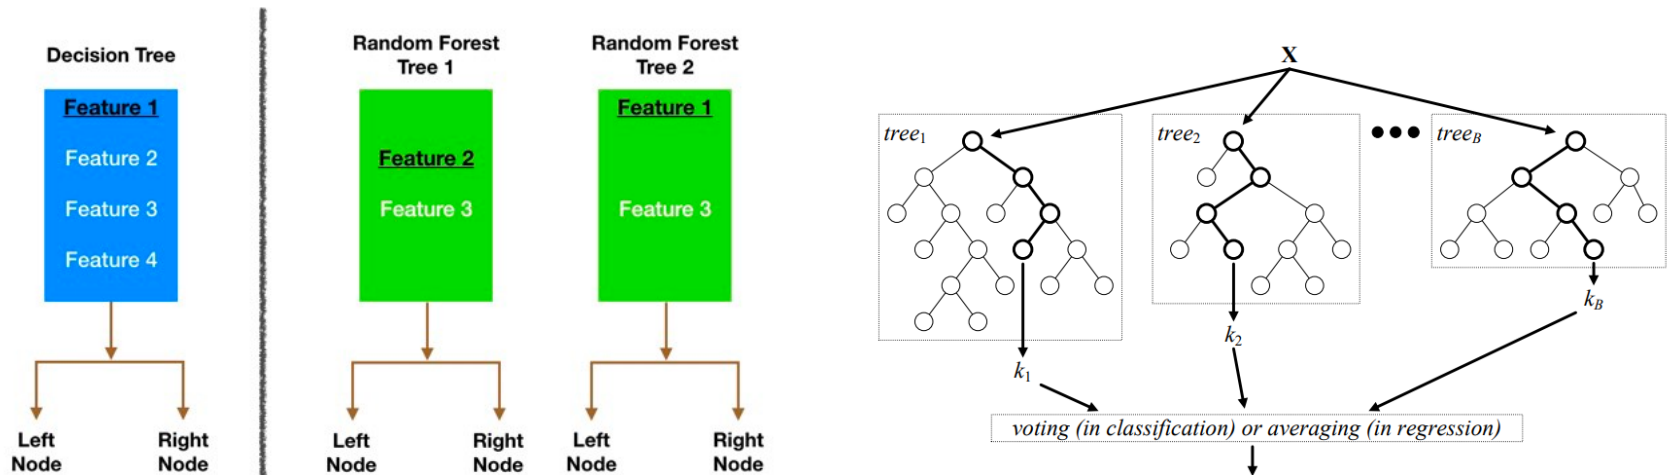
3. Output $\hat{f}(x) = f_M(x)$.

- **XGBoost** (eXtreme Gradient Boosting) : Faster and efficient implementation of Gradient Boosting

Random Forest

■ Random Forest

- Similar to Bagging, random forest consists of multiple trees (*forest*).
- Unlike bagging, it builds a large collection of *de-correlated* trees, and averages them.
- Generate different trees by using
 - Random *samples of training data points* when building trees
 - Random *subsets of features* when splitting nodes
- Random forests get their name from injecting randomness into the tree building to ensure each tree is different.



Random Forest

- Two ways in which the trees in a random forest are randomized
 - **by selecting the data points used to build a tree**
 - **bootstrap**: It leads to each decision tree in the random forest being built on a slightly different dataset.
 - From a list ['a', 'b', 'c', 'd'], possible examples of bootstrap samples are ['b', 'd', 'd', 'c'] and ['d', 'a', 'd', 'a'].
 - **by selecting the features in each split test.**
 - **max_features**: in each node, the algorithm randomly selects a subset of the features, and it looks for the best possible test involving one of these features
 - each node in a tree can make a decision using a different subset of the features.
 - A high **max_features** means that the trees in the random forest will be quite similar, and they will be able to fit the data easily, using the most distinctive features.
 - A low **max_features** means that the trees in the random forest will be quite different, and that each tree might need to be very deep in order to fit the data well.
- * Typically, for a classification problem with p features, \sqrt{p} (rounded down) features are used in each split. For regression problems the inventors recommend $p/3$ (rounded down) as the default. In practice the best value will depend on the problem.

Feature Importance

■ Feature Importance in Decision Tree

- Feature importance summarizes the workings of a tree by rating how important each feature is for the decision the tree makes.
- The importance of a feature is computed as the total reduction of the criterion brought by that feature. (Thus, it is called as impurity-based feature importance)
- It is a number between 0 and 1 for each feature, where 0 means “not used at all” and 1 means “perfectly predicts the target.”

■ Feature Importance in Ensemble models

- Similar to the decision tree, the ensemble models provide feature importances.
- Computed by aggregating the feature importances over the trees in the forest.
- Typically, the feature importances provided by the random forest are more reliable than the ones provided by a single tree.
- cf. https://scikit-learn.org/stable/auto_examples/inspection/plot_permutation_importance.html