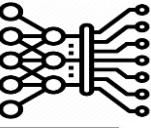


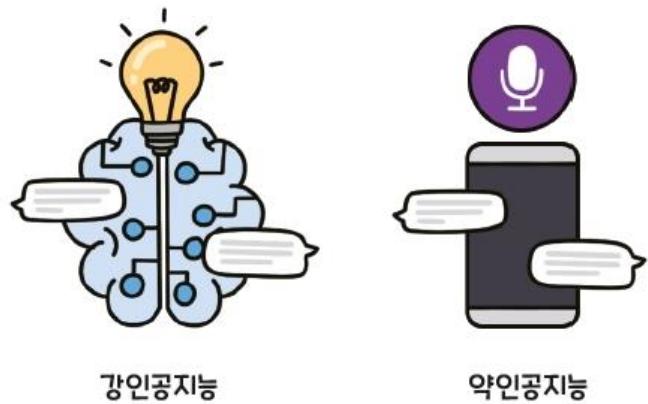
# **What is Machine Learning**

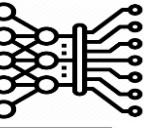
## **(Basic Machine Learning)**



## ❖ Concepts of AI, ML, and DL

- Artificial intelligence
  - A technology that embodies the intellectual abilities of humans through computers
- Classification of artificial intelligence
  - Strong AI: AI with performance beyond human capabilities
  - Weak AI: AI designed for use as a tool in certain areas

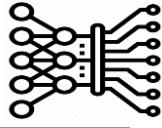




## ❖ Concepts of AI, ML, and DL

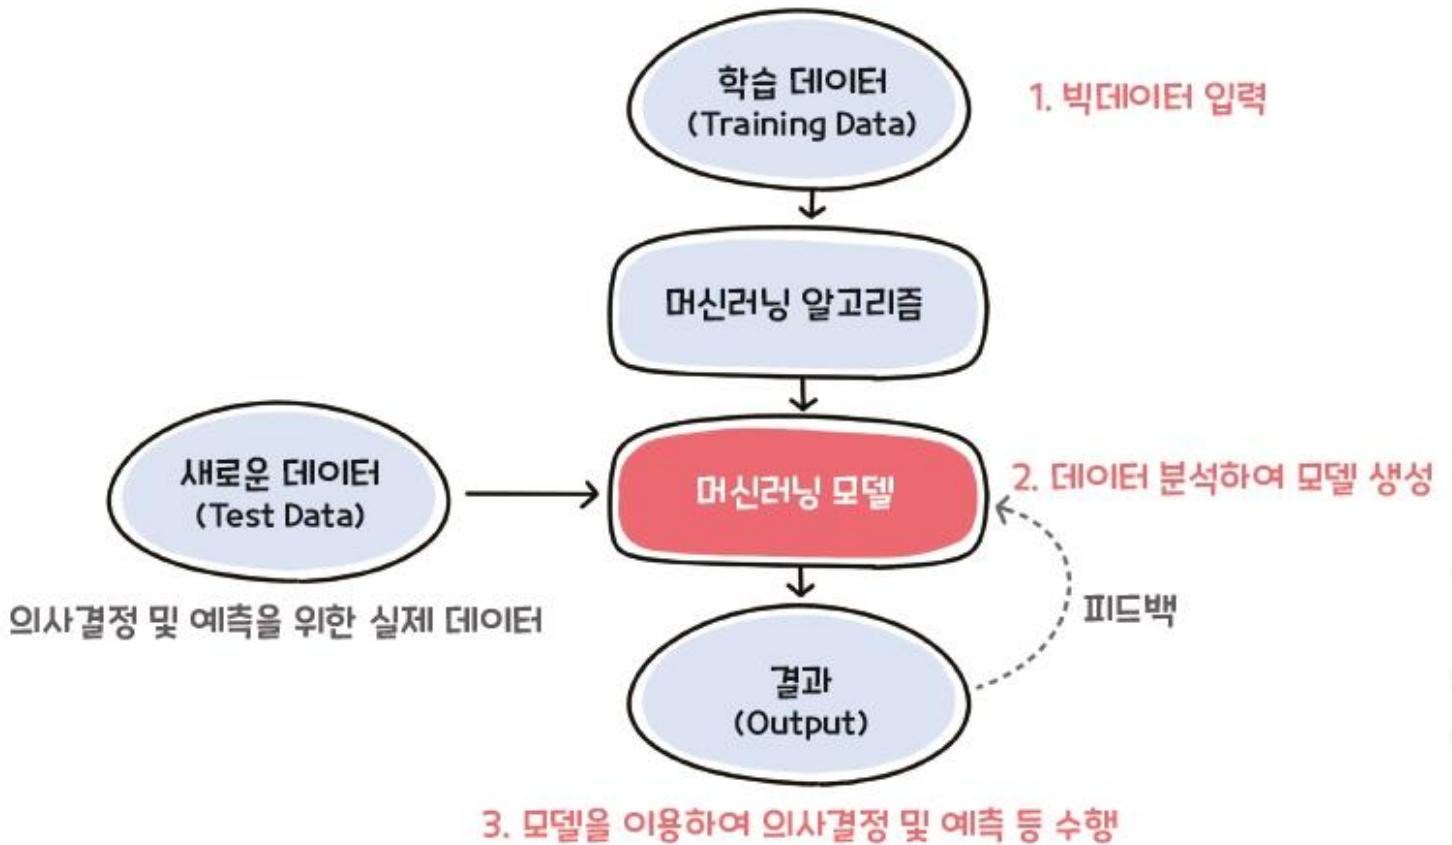
### ▪ Machine Learning

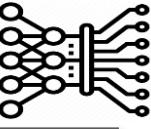
- A technology that allows computers to learn like humans so that computers themselves can discover new rules without human help
- Machine learning basically analyzes data using algorithms, learns through analysis, and makes judgments or predictions based on what is learned
- Machine learning is the process of self-learning and processing data
  - Insert Big Data
  - Analyze data to create a model
  - Use models to make decisions, predictions, etc



## ❖ Concepts of AI, ML, and DL

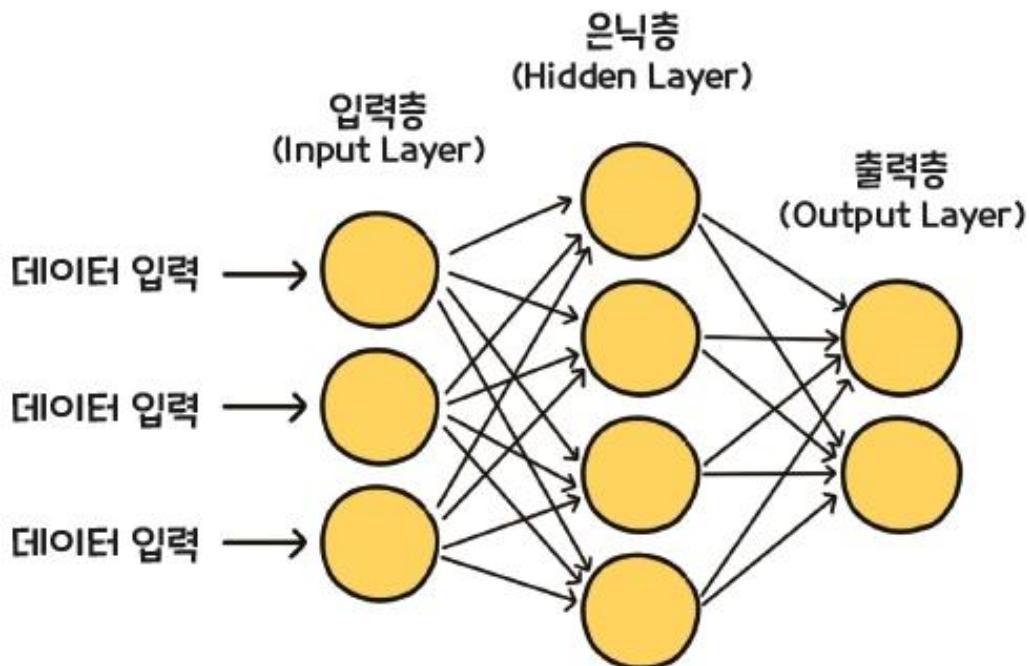
### ▪ Machine Learning

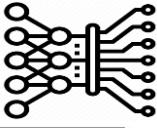




## ❖ Concepts of AI, ML, and DL

- Deep Learning
  - ANN, Artificial Neural Network
    - A network of interconnected neurons

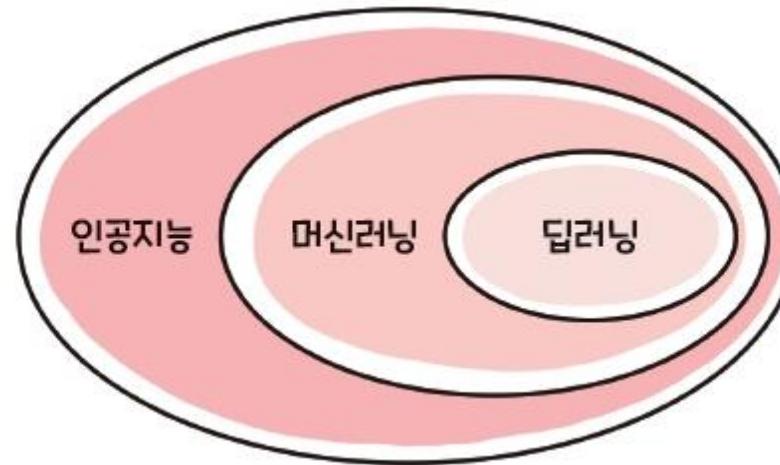


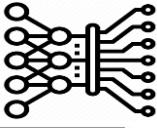


## ❖ Concepts of AI, ML, and DL

### ▪ Deep Learning

- Technology for performing machine learning using artificial neural networks with multiple hidden layers
- “Deep” in deep learning means deep layers of continuous neural networks
- Performance increases as this neural network deepens

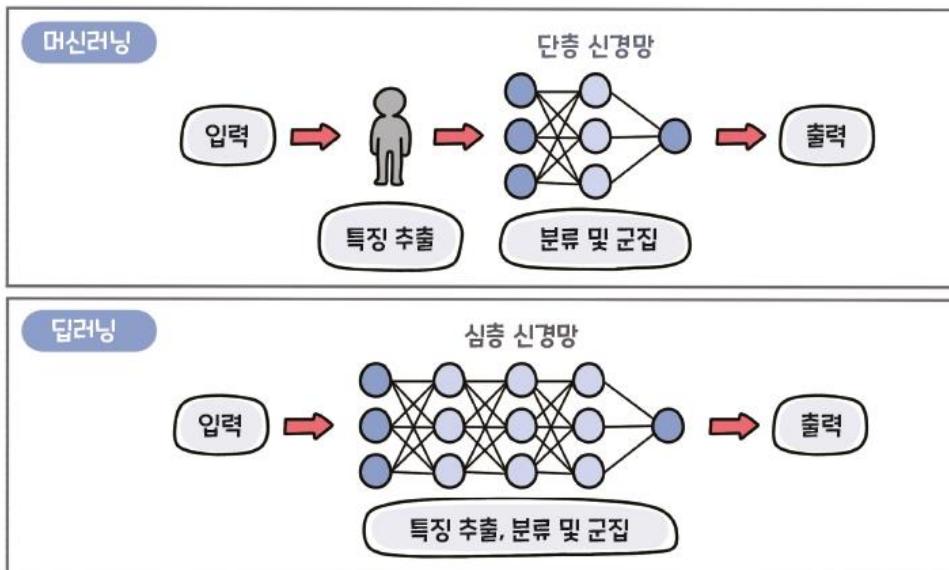


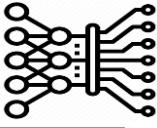


## ❖ Difference between ML and DL

### ▪ Human in the loop

- Machine learning involves some degree of intervention, such as human informing the learning data of labels (corrects) or extracting the characteristics of the data
- Deep learning learns on its own without human intervention



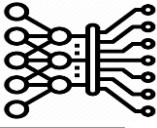


## ❖ Difference between ML and DL

### ▪ Feature extraction

- In machine learning, in order for a computer to learn on its own, it has to convert human-recognized data into computer-aware data
- For this task, it finds out what characteristics each data has and converts the data into vector

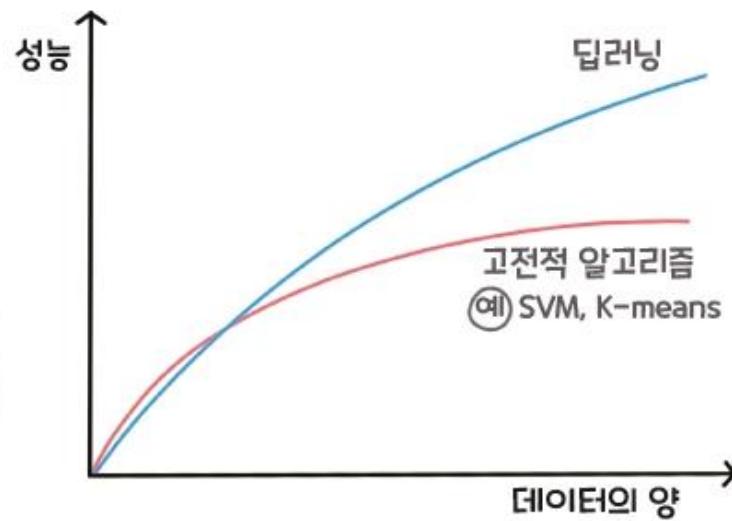


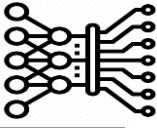


## ❖ Difference between ML and DL

### ▪ Data dependencies

- Deep learning directly extracts important features to solve a given problem
- If you don't have enough data, you can't extract the exact features
- On the other hand, if sufficient data is given, it performs well enough to identify important features that humans do not recognize

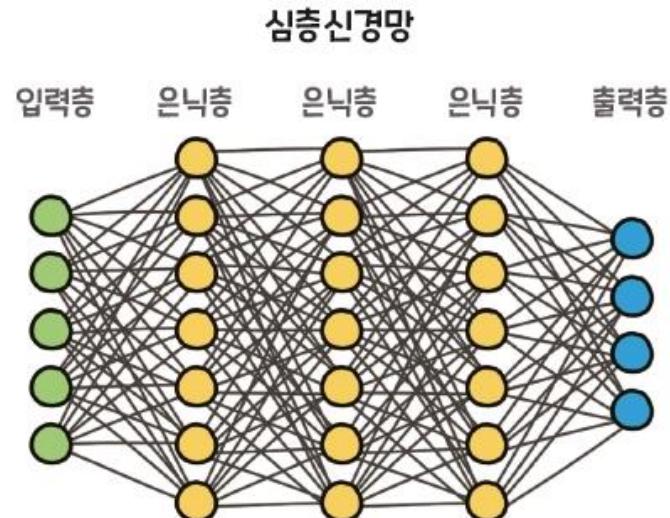
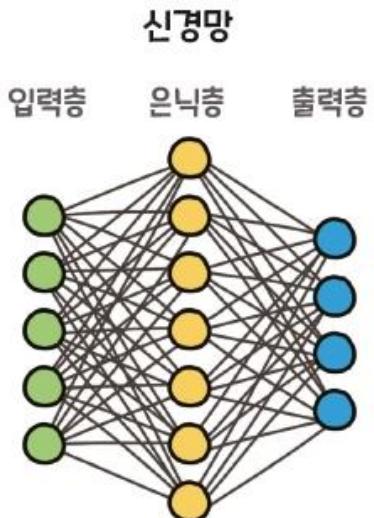


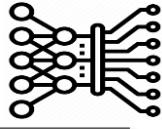


## ❖ Difference between ML and DL

- Using neural network

- Deep learning uses a deep neural network to extract features from input data and derive results (prediction or classification) on its own
- The use of deep neural networks is a distinct characteristic of deep learning



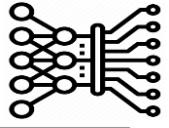


## ❖ Difference between ML and DL

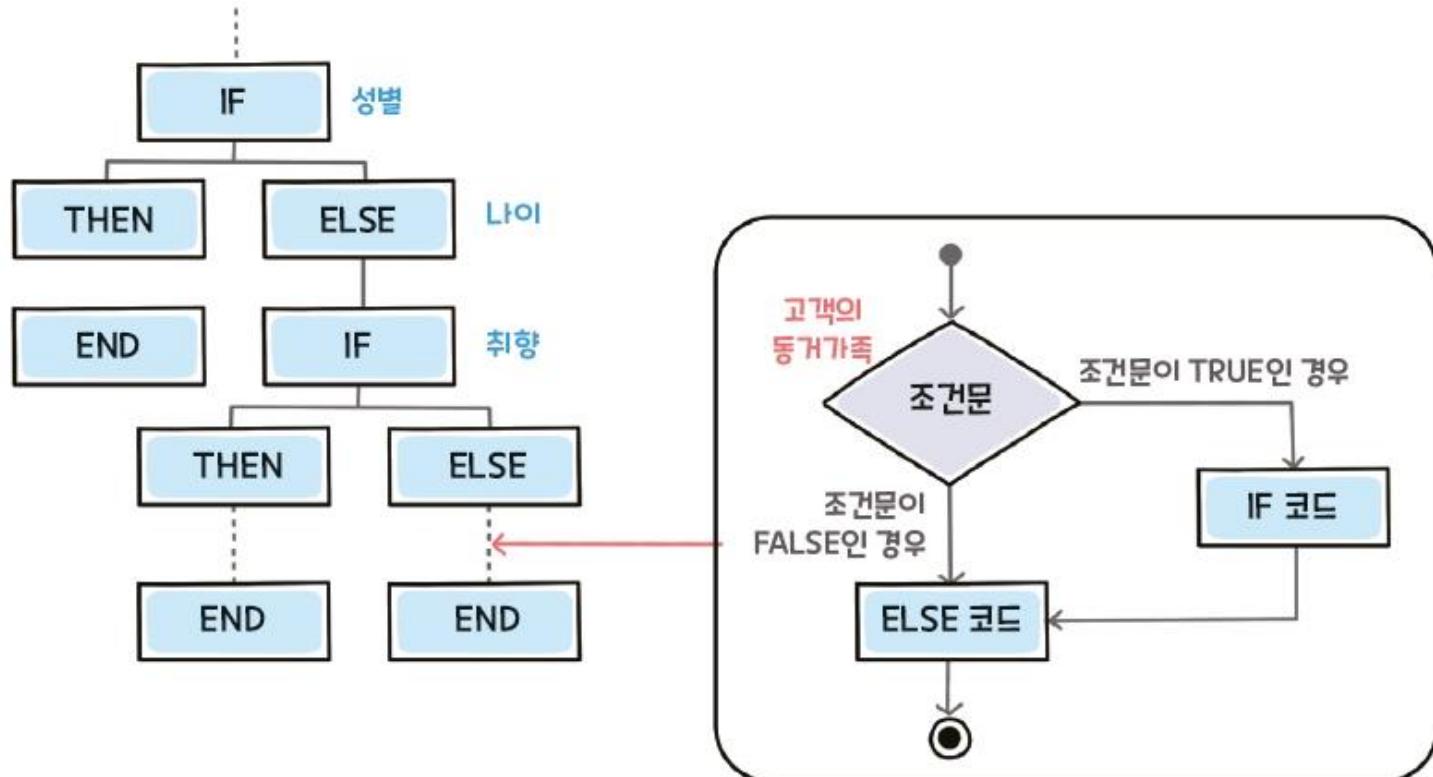
- Using neural network

구분	머신러닝	딥러닝
필요한 데이터의 양	적은 양의 데이터도 가능	빅데이터
정확도	낮음	높음
훈련 시간	짧은 시간 안에 가능	오래 걸림
하드웨어	CPU만으로도 가능	GPU
하이퍼파라미터 튜닝	제한적	다양한 방법으로 튜닝 가능

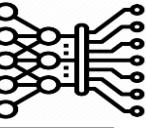
# Why Machine Learning?



## ❖ Limitation of basic programming

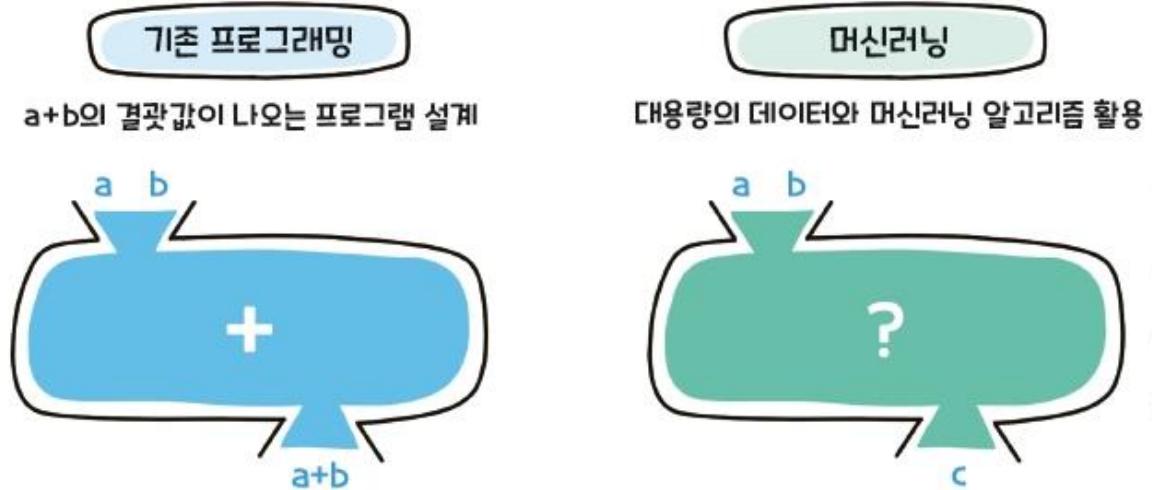


# Why Machine Learning?

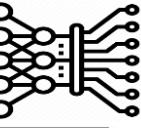


## ❖ Usability of machine learning

- But it's not the right time to make a quick decision
  - Using machine learning to solve this problem
- Machine learning is a very useful solution when large amounts of data and many variables are involved, and programs with conventional rules cannot solve complex tasks or problems



# Kinds of Machine Learning

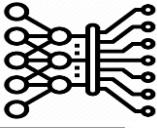


## ❖ Categorization for training

- Supervised learning : classification and regression
- Unsupervised learning : clustering
- Reinforcement learning : use rewards for actions taken in the environment to conduct learning

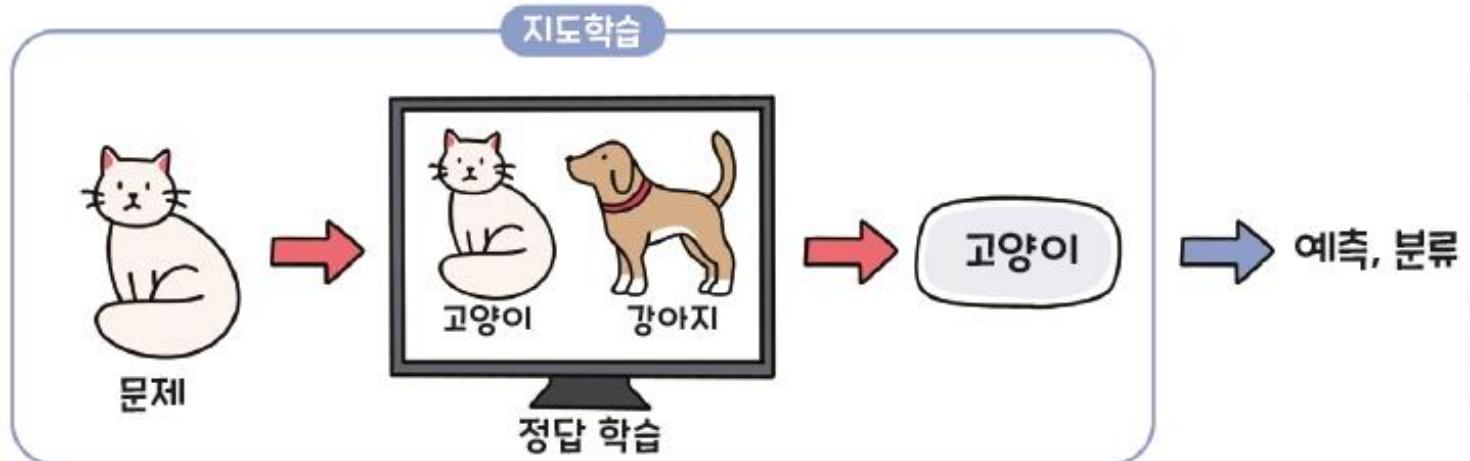


# Kinds of Machine Learning

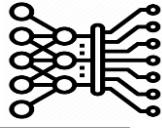


## ❖ Supervised learning

- Learning to predict the right answer to an unknown problem by learning questions and answers together
- The models used in supervised learning include prediction and classification

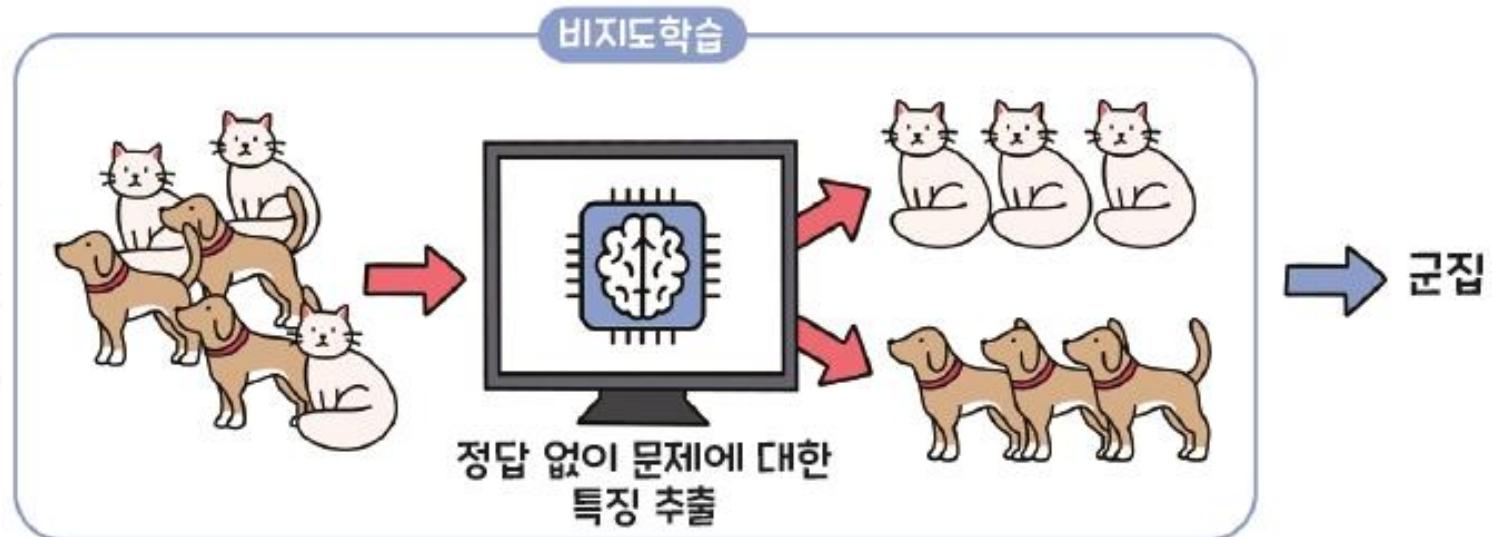


# Kinds of Machine Learning

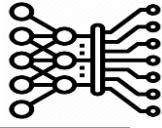


## ❖ Unsupervised learning

- A form of computer learning without the help
- Computer uses training data to find regularity between data



# Kinds of Machine Learning

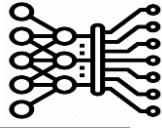


## ❖ Unsupervised learning

- Unlike supervised learning, which identified the relationship between x (input data) and y (labels in supervised learning),
- Unsupervised learning identifies the relationship between x by itself
- In other words, the difference between y (label)
  - Clustering is a model used in unsupervised learning

구분	지도학습	비지도학습
필요한 데이터 종류	x(학습 데이터), y(레이블)	x(학습 데이터)

# Kinds of Machine Learning



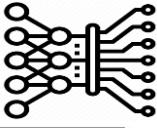
## ❖ Reinforcement learning

- Learning to be rewarded for what you've done
- How computers learn to choose the best behavior for a given state



# Kinds of Machine Learning

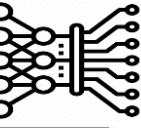
---



## ❖ Reinforcement learning

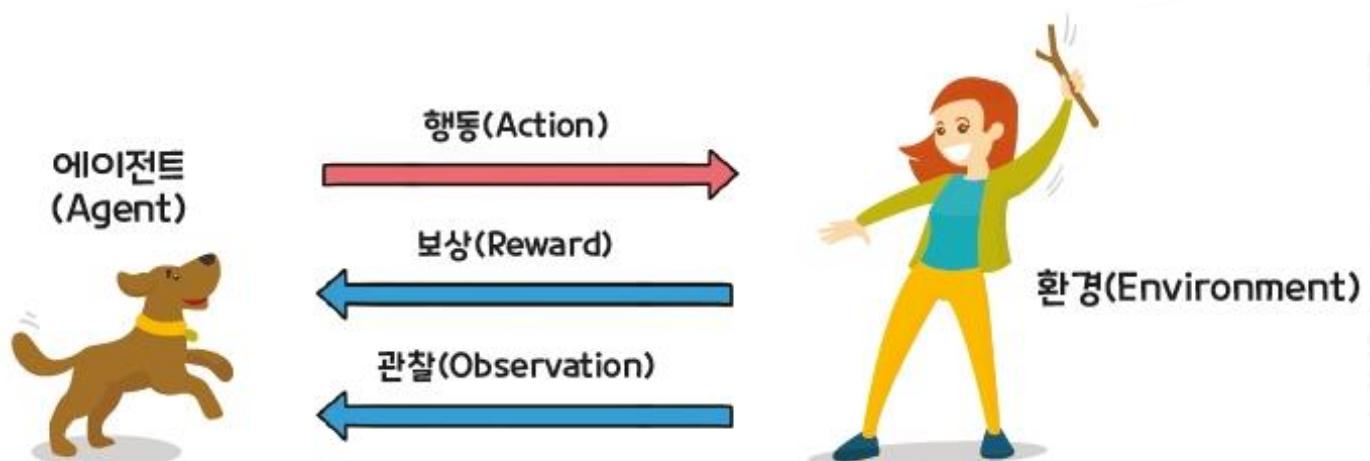
- Agent: Subject to act in a given problem situation
- State : Current situation
- Action: Options that the player can take
- Rewards: Benefits that follow when a player does something
- Environment: means the problem itself
- Observation : Information about the collected by the agent

# Kinds of Machine Learning

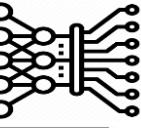


## ❖ Reinforcement learning

- Depending on the behavior chosen by the agent in a given environment, you are rewarded if the behavior is the right choice, and punished if the behavior is the wrong choice
- Reinforcement learning allows the agent to keep an eye on the status and learn (behavior) toward higher rewards

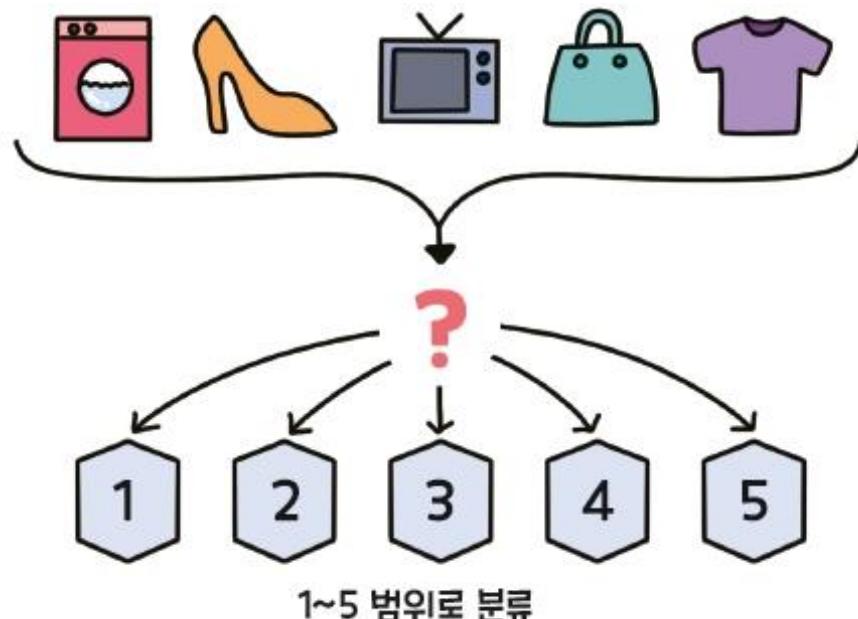


# Types of Machine Learning Algorithm

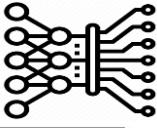


## ❖ Classification

- A technique for learning labeled data, classifying data with similar properties, and finding out which group the newly entered data

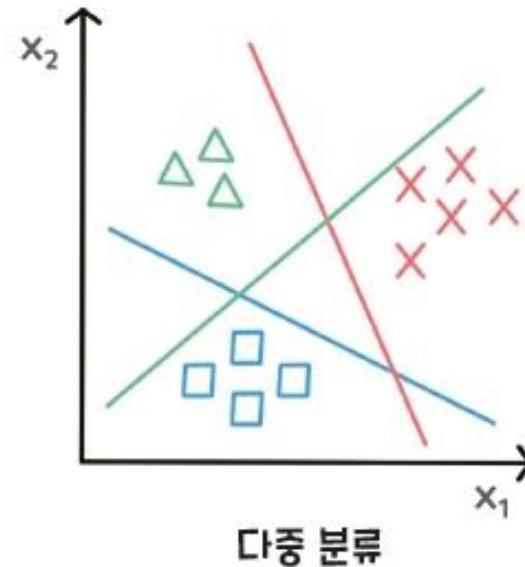
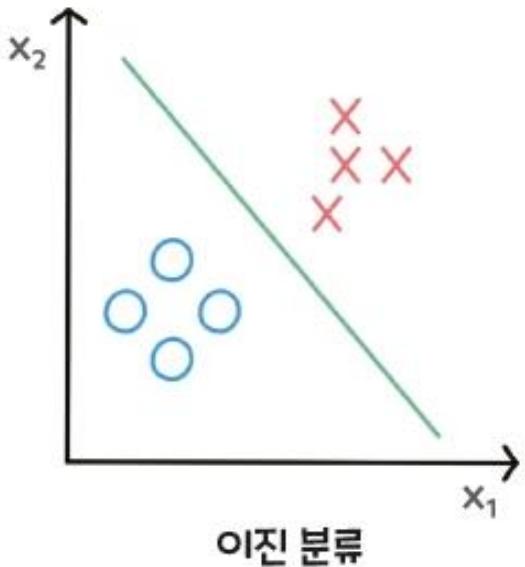


# Types of Machine Learning Algorithm



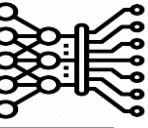
## ❖ Classification

- Binary classification : categorize data into 2 groups
- Multiclass classification : categorize data into 3 or more groups



# Types of Machine Learning Algorithm

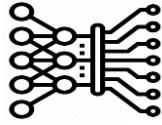
---



## ❖ Classification (Algorithm)

- K-neighbor nearest
- Support vector machine
- Decision tree

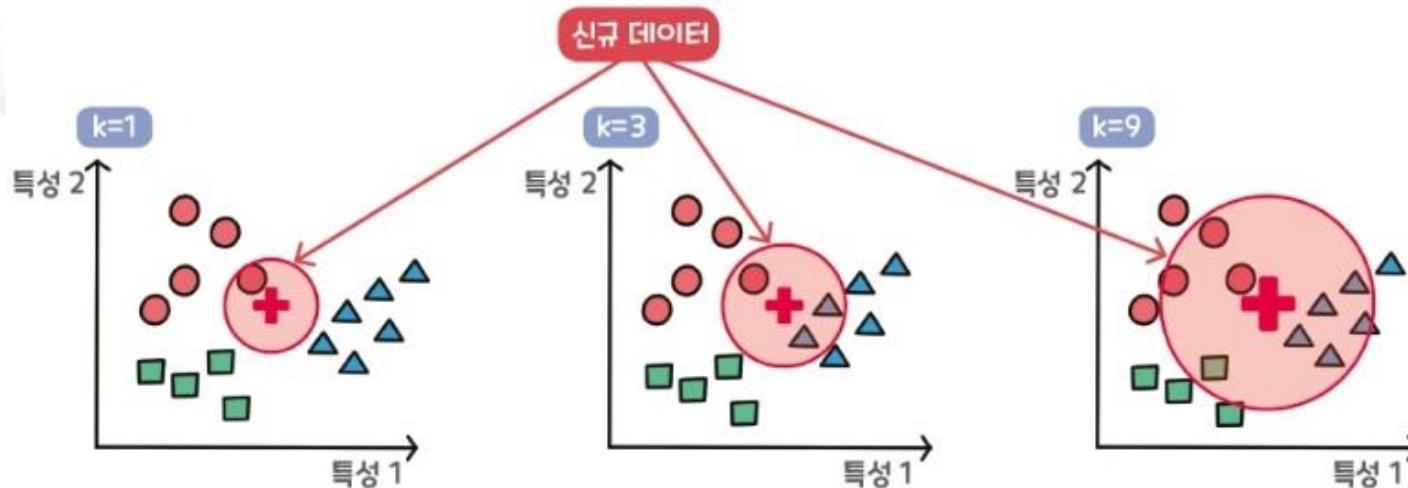
# Types of Machine Learning Algorithm



## ❖ Classification (Algorithm)

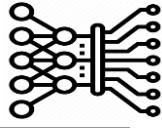
### ▪ K-neighbor nearest

- Algorithms to classify which of the existing groups of data (K groups) belongs to when new data comes in
- (Example) When new data is entered when  $K=1$ , new data is classified as a red circle, when  $K=3$ , and when  $K=9$ , it is classified as a blue triangle



# Types of Machine Learning Algorithm

---

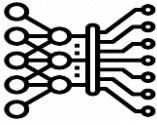


## ❖ Classification (Algorithm)

### ▪ K-neighbor nearest

- KNNs are not significantly affected by the noise present in the learning data and are quite effective when the number of learning data is large
- However, it is unclear which hyperparameters are suitable for analysis, so there is a disadvantage that researchers should randomly select according to each characteristic of the data

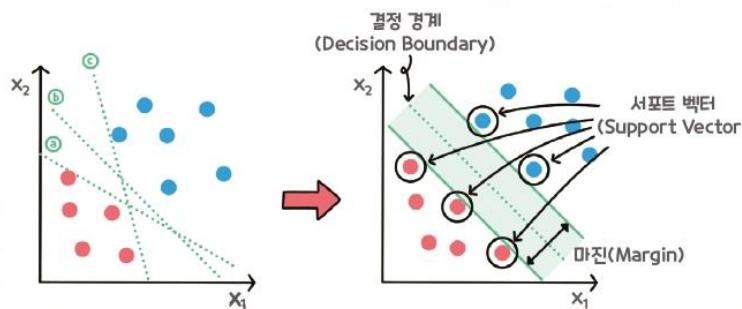
# Types of Machine Learning Algorithm



## ❖ Classification (Algorithm)

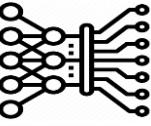
### ▪ Support vector machine

- Categorize data in the direction of maximizing margin, which means margin between two categories
- SVMs find and classify lines that maximize margins, so larger margins are more likely to be classified even if new data comes in
- SVM is easy to use and highly predictive
  - However, it takes time to build a model and the results are less descriptive



- 결정 경계(Decision Boundary) : 분류를 위한 기준선
- 서포트 벡터(Support Vector) : 결정 경계와 가장 가까운 위치에 있는 데이터
- 마진(Margin) : 결정 경계와 서포트 벡터 사이의 거리

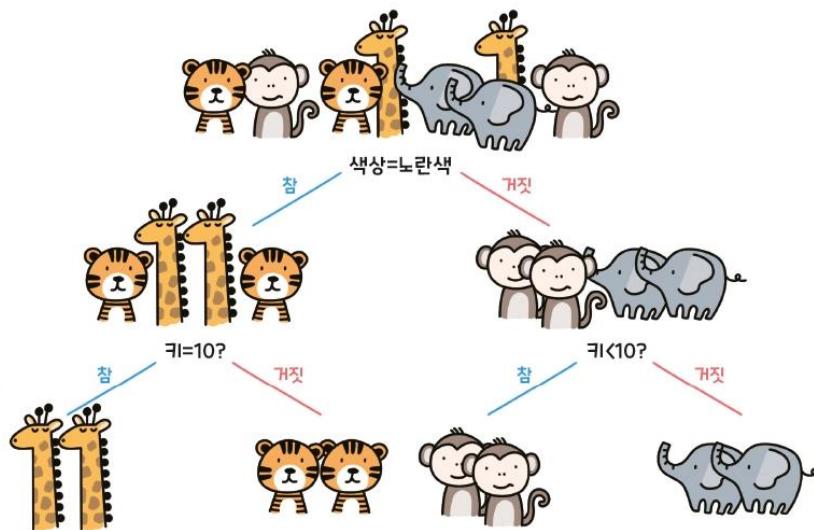
# Types of Machine Learning Algorithm



## ❖ Classification (Algorithm)

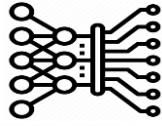
### ▪ Decision tree

- An analysis method for classifying decision-making rules into tree forms
- It is called 'decision tree' because the method of starting from the upper node and expanding to the lower node according to the classification criteria resembles 'tree'



# Types of Machine Learning Algorithm

---

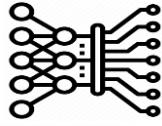


## ❖ Classification (Algorithm)

### ▪ Decision tree

- Decision Tree is intuitive and easy to understand the analysis process
- In the case of artificial neural networks, it is a black box model that is difficult to explain the analysis results, while decision trees can observe the analysis process with their eyes
- Need for a clear explanation of the results

# Types of Machine Learning Algorithm



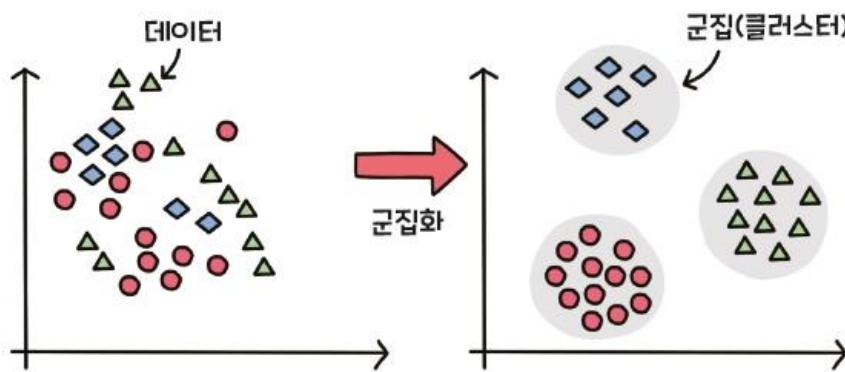
## ❖ Unsupervised learning

### ▪ Cluster

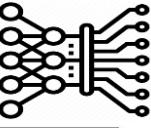
- A group of data with similar characteristics

### ▪ Clustering

- Classifying the data into clusters according to a similar degree when given the data
- Various data are mixed together, but the clustering process groups similar data as shown in the graph on the right



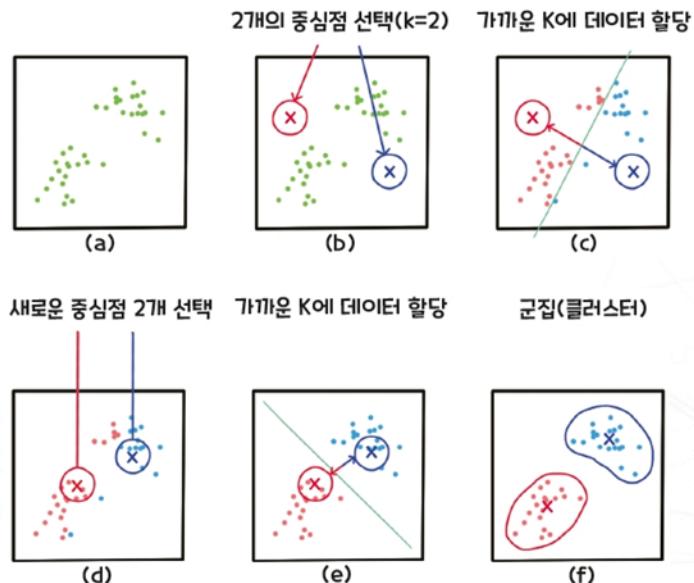
# Types of Machine Learning Algorithm



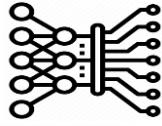
## ❖ Unsupervised learning

### ▪ K-means clustering

- 'K' is the number of groups to be grouped from the given data
- 'Means' means the average distance between the center of each cluster and the data
- The center of the cluster is called centroids



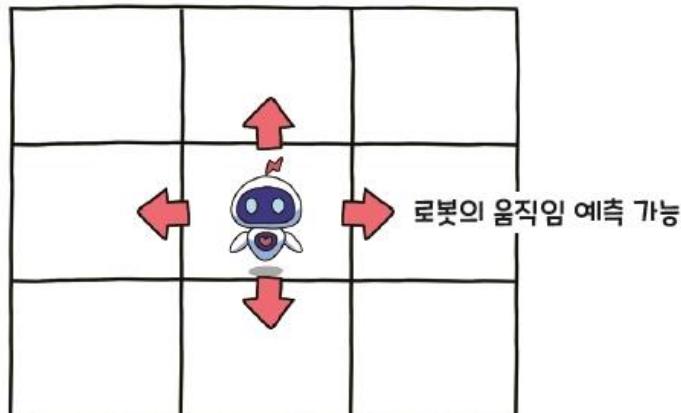
# Types of Machine Learning Algorithm



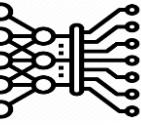
## ❖ Reinforcement learning

### ▪ Algorithm

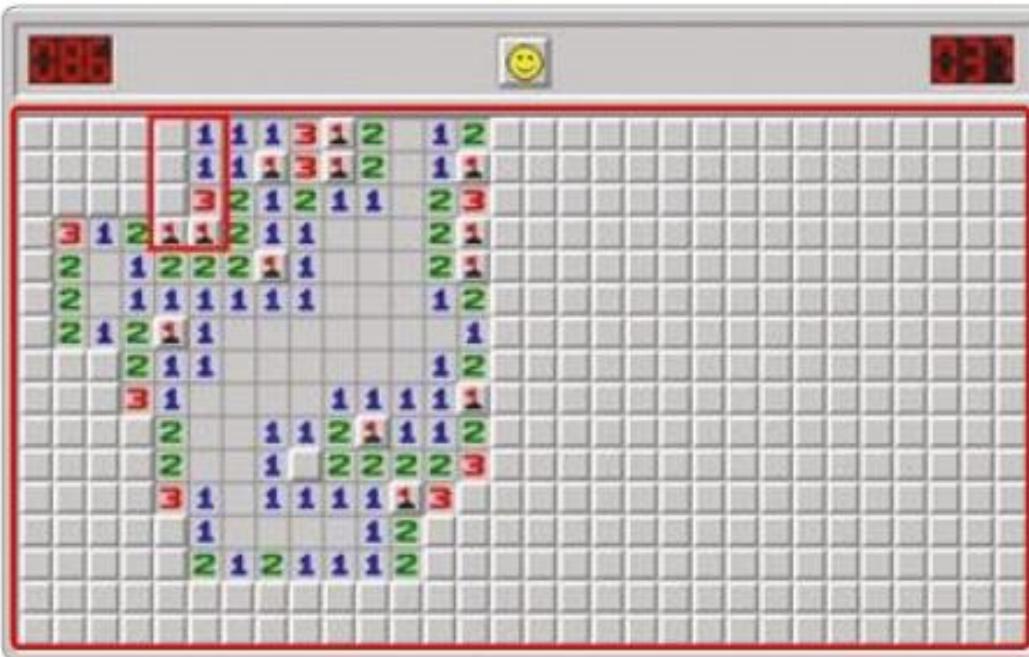
- Model-based algorithms refer to the probability that an action in the current state will result in the next state
- Intuitive visibility of the robot's next state as it moves up, down, left, and right in a grid space
  - Model-based algorithms can predict changes in state according to behavior, resulting in optimal solutions

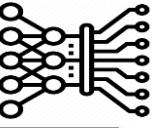


# Types of Machine Learning Algorithm



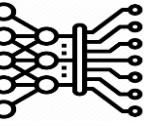
- ❖ Reinforcement learning
  - Algorithm
    - Finding a policy that maximizes the rewards an agent receives through action





- ❖ Install the scikit-learn library
  - pip install scikit-learn

**TIP** [https://scikit-learn.org/dev/\\_downloads/scikit-learn-docs.pdf](https://scikit-learn.org/dev/_downloads/scikit-learn-docs.pdf)에 접속하면 가장 최신 버전의 사이킷 런 사용 설명서를 무료로 다운로드할 수 있다. 무려 2,500여 쪽에 달하는 방대한 문서다. 그렇다고 겁먹을 필요는 없다. 필요한 부분을 선택적으로 참조하면 된다.



## ❖ Load the dataset

프로그램 3-1(a) iris 데이터셋 읽기

```
01 from sklearn import datasets  
02  
03 d=datasets.load_iris()      # iris 데이터셋을 읽고  
04 print(d.DESCR)            # 내용을 출력
```

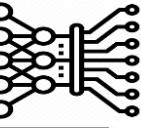
- 01행: sklearn 모듈의 datasets 클래스를 불러옴
- 03행: load\_iris 함수를 호출해 iris 데이터셋을 읽어 객체 d에 저장
- 04행: 객체 d의 DESCR 변수를 출력

## ❖ Terminology

- Dataset
- Feature vector
- Class

**TIP** 기계 학습이 사용하는 데이터는 여러 개의 샘플을 담고 있어서 데이터셋(data set)이라 부르기도 한다. 이 책에서는 데이터와 데이터셋을 엄밀히 구분하지 않고 함께 사용하는데, 데이터셋은 iris처럼 특정한 데이터를 가리킬 때 주로 사용한다.

# Load 'iris' Dataset



## Iris plants dataset

### \*\*Data Set Characteristics:\*\*

:Number of Instances: 150 (50 in each of three classes)

:Number of Attributes: 4 numeric, predictive attributes and the class

:Attribute Information:

- sepal length in cm

- sepal width in cm

- petal length in cm

- petal width in cm

- class:

- Iris-Setosa

- Iris-Versicolour

- Iris-Virginica

:Summary Statistics:

	Min	Max	Mean	SD	Class Correlation
sepal length:	4.3	7.9	5.84	0.83	0.7826
sepal width:	2.0	4.4	3.05	0.43	-0.4194
petal length:	1.0	6.9	3.76	1.76	0.9490 (high!)
petal width:	0.1	2.5	1.20	0.76	0.9565 (high!)

:Missing Attribute Values: None

:Class Distribution: 33.3% for each of 3 classes.

:Creator: R.A. Fisher

:Donor: Michael Marshall (MARSHALL%PLU@io.arc.nasa.gov)

:Date: July, 1988

...

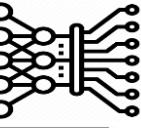
150개의 샘플

네 개의 특징(feature)

세 개의 부류



# Load 'iris' Dataset



## ❖ 'iris' dataset

프로그램 3-1(b) iris의 내용 살펴보기

```
05 for i in range(0,len(d.data)):      # 샘플을 순서대로 출력  
06     print(i+1,d.data[i],d.target[i])
```

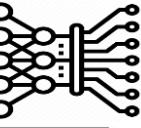
```
1 [5.1 3.5 1.4 0.2] 0  
2 [4.9 3.  1.4 0.2] 0  
3 [4.7 3.2 1.3 0.2] 0  
4 [4.6 3.1 1.5 0.2] 0  
...  
...
```

```
51 [7.  3.2 4.7 1.4] 1  
52 [6.4 3.2 4.5 1.5] 1  
53 [6.9 3.1 4.9 1.5] 1  
54 [5.5 2.3 4.  1.3] 1  
...  
101 [6.3 3.3 6.  2.5] 2  
102 [5.8 2.7 5.1 1.9] 2  
103 [7.1 3.  5.9 2.1] 2  
104 [6.3 2.9 5.6 1.8] 2  
...  
...
```

d.data(특징 벡터)

d.target(레이블)

# Representation of dataset

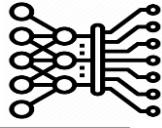


- ❖ Representing samples as feature vectors and labels

- Feature vectors are denoted by  $\mathbf{x}$  특징 벡터:  $\mathbf{x}=(x_1, x_2, \dots, x_d)$ 
    - $d$  is the number of features called the dimension of the feature vector
  - Labels are 0,1,2,...A value of , $c-1$  or 1,2,...A value of , $c-1,c$  or one hot code
    - One hot code is a binary sequence with only one element
    - Ex) Setosa: (1,0,0), Versicolor: (0,1,0), Virginica: (0,0,1)

특징 벡터 $\mathbf{x} = (x_1, x_2, \dots, x_d)$	레이블(참값) $y$
샘플 1: (5.1, 3.5, 1.4, 0.2)	0
샘플 2: (4.9, 3.0, 1.4, 0.2)	0
...	...
샘플 51: (7.0, 3.2, 4.7, 1.4)	1
샘플 52: (6.4, 3.2, 4.5, 1.5)	1
...	...
샘플 101: (6.3, 3.3, 6.0, 2.5)	2
샘플 102: (5.8, 2.7, 5.1, 1.9)	2
...	...
샘플 n: (5.9, 3.0, 5.1, 1.8)	2

# Data Distribution of Feature Space



## ❖ iris dataset

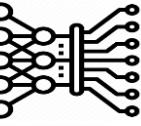
- Distribution of data in a three-dimensional space, excluding one data dimension

프로그램 3-2

iris 데이터의 분포를 특징 공간에 그리기

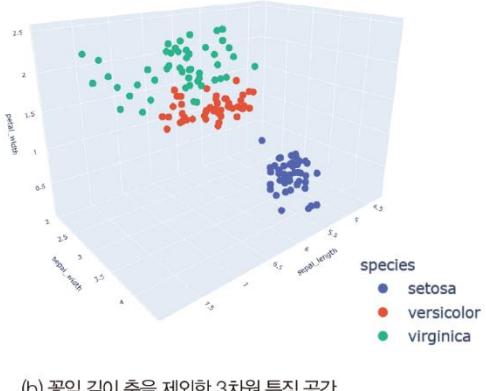
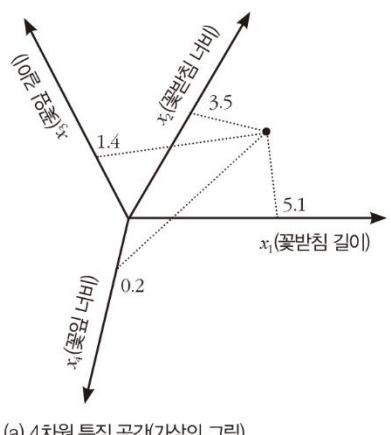
```
01 import plotly.express as px  
02  
03 df = px.data.iris()  
04 fig = px.scatter_3d(df, x='sepal_length', y='sepal_width', z='petal_width',  
color='species')      # petal_length를 제외하여 3차원 공간 구성  
05 fig.show(renderer="browser")
```

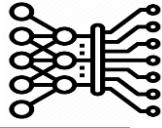
# Data Distribution of Feature Space



## ❖ Observe the distribution of data in the feature space

- Setosa is distributed downward and Virginica is distributed upward for the vertical width
  - Petal width is excellent in discernment
- The segmental width axis overlaps a lot in three categories, so it is less sensible
- As a whole, the three classes occupy different areas of the three-dimensional space, with several samples overlapping

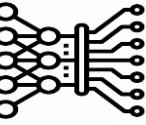




## NOTE 다차원 특징 공간

종이에 그릴 수 있는 공간은 3차원으로 제한되지만, 수학은 아주 높은 차원까지 다룰 수 있다. 예를 들어 2차원 상의 두 점  $\mathbf{x} = (x_1, x_2)$ 과  $\mathbf{y} = (y_1, y_2)$ 의 거리를  $d(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2}$ 으로 계산할 수 있는데, 4차원 상의 두 점  $\mathbf{x} = (x_1, x_2, x_3, x_4)$ 과  $\mathbf{y} = (y_1, y_2, y_3, y_4)$ 의 거리는  $d(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2 + (x_3 - y_3)^2 + (x_4 - y_4)^2}$ 로 계산 할 수 있다.

일반적으로  $d$ 차원 상의 두 점의 거리는  $d(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{i=1}^d (x_i - y_i)^2}$ 로 계산한다. 기계 학습에서는  $d =$ 수백~수만에 달 하는 매우 고차원 특징 공간의 데이터를 주로 다룬다.



## ❖ Using the support vector machine model

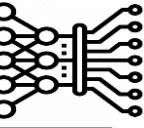
프로그램 3-1(c)

iris에 기계 학습 적용: 모델링과 예측

```
07 from sklearn import svm           Hyperparameter
08
09 s=svm.SVC(gamma=0.1,C=10)          # SVM 분류 모델 SVC 객체 생성하고
10 s.fit(d.data,d.target)             # iris 데이터로 학습
11
12 new_d=[[6.4,3.2,6.0,2.5],[7.1,3.1,4.7,1.35]]  # 101번째와 51번째 샘플을 변형하여
13 res=s.predict(new_d)               # 새로운 데이터 생성
14 print("새로운 2개 샘플의 부류는", res)
```

새로운 2개 샘플의 부류는 [2 1]

- 09행: SVM의 분류기 모델 SVC 클래스의 객체를 생성하여 s에 저장
- 10행: 객체 s의 fit 함수는 훈련 집합을 가지고 학습을 수행  
(매개변수로 특징 벡터 iris.data와 레이블 iris,target을 설정)
- 13행: 객체 s의 predict 함수는 테스트 집합을 가지고 예측 수행



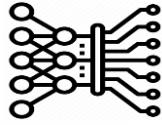
## ❖ The importance of objective performance measurements

- Important when choosing a model
- Important when deciding whether to install on-site

## ❖ Generalization capabilities

- Performance on new data not used for learning
- The most obvious way is to install it on-site and measure performance
- Cost makes it difficult to apply it in real life
- Requires wisdom to segment and use given data

# Performance Measurement



## ❖ Confusion matrix

- Matrix recording the number of correct and incorrect classifications by class
  - $n_{ij}$ 는 모델이  $i$ 라고 예측했는데 실제 부류는  $j$ 인 샘플의 개수

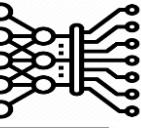
		참값(그라운드 트루스)					
		부류 1	부류 2	...	부류 $j$	...	부류 $c$
예측한 부류	부류 1	$n_{11}$	$n_{12}$		$n_{1j}$		$n_{1c}$
	부류 2	$n_{21}$	$n_{22}$		$n_{2j}$		$n_{2c}$
	...						
	부류 $i$	$n_{i1}$	$n_{i2}$		$n_{ij}$		$n_{ic}$
	...						
	부류 $c$	$n_{c1}$	$n_{c2}$		$n_{cj}$		$n_{cc}$

(a) 부류가  $c$ 개인 경우

		그라운드 트루스	
		긍정	부정
예측값	긍정	TP	FP
	부정	FN	TN

(b) 부류가 2개인 경우

- Positive and negative negative in binary classification
- True positive, false negative, false positive, true negative



## ❖ Performance metric

- Accuracy

$$\text{정확률} = \frac{\text{맞힌 샘플 수}}{\text{전체 샘플 수}} = \frac{\text{대각선 샘플 수}}{\text{전체 샘플 수}}$$

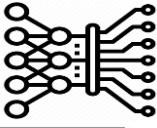
- Specificity and sensitivity

$$\text{특이도} = \frac{\text{TN}}{\text{TN} + \text{FP}}, \text{ 민감도} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

- Precision and recall

$$\text{정밀도} = \frac{\text{TP}}{\text{TP} + \text{FP}}, \text{ 재현율} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

# Divide into Training/Validation/Test



## ❖ Training/Validation/Test

- Training set

- Data used to learn machine learning models that provide both feature vector and label information

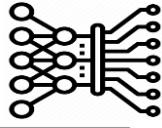
- Test Set

- Data used to measure the performance of a learned model, which provides only feature vector information when predicting, and uses label information when measuring accuracy with prediction results

### NOTE 하이퍼 매개변수 설정

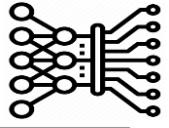
하이퍼 매개변수<sup>hyper parameter</sup>란 모델의 동작을 제어하는 데 쓰는 변수이다. 모델의 학습을 시작하기 전에 설정해야 하는데, 적절한 값으로 설정해야 좋은 성능을 얻을 수 있다. 최적의 하이퍼 매개변수 값을 자동으로 설정하는 일을 하이퍼 매개변수 최적화(hyper parameter optimization)라하는데, 이것은 기계 학습의 중요한 주제 중 하나다. 하이퍼 매개변수 최적화는 4.10절에서 다룬다.

# Divide into Training/Validation/Test



- ❖ Divide the given data into training, validation, and test sets at an appropriate rate
  - Model selection included: divided into training/validation/test sets
  - Exclude model selection: split into training/test sets





## ❖ Exclude the model selection

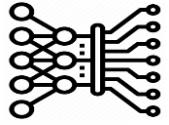
- 08행: train\_test\_split 함수로 훈련 60%, 테스트 40%로 랜덤 분할
- 12행: 훈련 집합 x\_train, y\_train을 fit 함수에 주어 학습 수행
- 14행: 테스트 집합의 특징 벡터 x\_test를 predict 함수에 주어 예측 수행
- 17~20행: 테스트 집합의 레이블 y\_test를 가지고 혼동 행렬 계산

프로그램 3-5

필기 숫자 인식 – 훈련 집합으로 학습하고 테스트 집합으로 성능 측정

```
01 from sklearn import datasets
02 from sklearn import svm
03 from sklearn.model_selection import train_test_split
04 import numpy as np
05
06 # 데이터셋을 읽고 훈련 집합과 테스트 집합으로 분할
07 digit=datasets.load_digits()
08 x_train,x_test,y_train,y_test=train_test_split(digit.data,digit.target,train_size=0.6)
09
```

# Divide into Training/Validation/Test

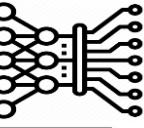


```
10 # svm의 분류 모델 SVC를 학습
11 s=svm.SVC(gamma=0.001)
12 s.fit(x_train,y_train)
13
14 res=s.predict(x_test)
15
16 # 혼동 행렬 구함
17 conf=np.zeros((10,10))
18 for i in range(len(res)):
19     conf[res[i]][y_test[i]]+=1
20 print(conf)
21
22 # 정확률 측정하고 출력
23 no_correct=0
24 for i in range(10):
25     no_correct+=conf[i][i]
26 accuracy=no_correct/len(res)
27 print("테스트 집합에 대한 정확률은", accuracy*100, "%입니다.")
```

예) 부류 3에 속하는 75개 샘플 중 73개를 3, 1개를 2, 1개를 7로 인식

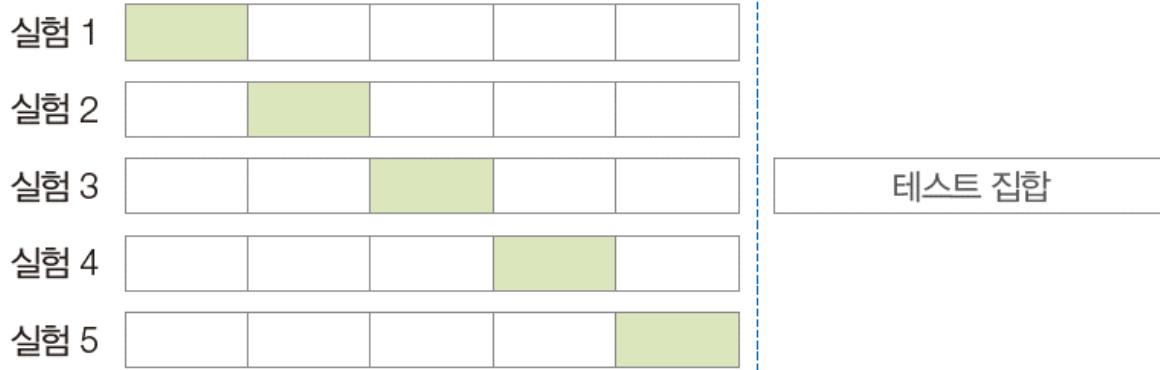
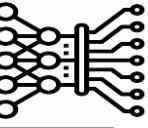
```
[[76.  0.  0.  0.  0.  0.  0.  0.  0.  0.]
 [ 0.  78.  0.  0.  0.  0.  0.  0.  3.  0.]
 [ 0.  0.  66.  1.  0.  0.  0.  0.  0.  0.]
 [ 0.  0.  0.  73.  0.  0.  0.  0.  0.  0.]
 [ 0.  0.  0.  0.  63.  0.  0.  0.  0.  0.]
 [ 0.  0.  0.  0.  0.  70.  0.  0.  0.  2.]
 [ 0.  0.  0.  0.  0.  0.  77.  0.  0.  0.]
 [ 0.  0.  0.  1.  0.  0.  0.  77.  0.  1.]
 [ 0.  0.  0.  0.  0.  0.  0.  0.  74.  0.]
 [ 0.  0.  0.  0.  0.  1.  0.  0.  0.  56.]]
```

테스트 집합에 대한 정확률은 98.74826147426981%입니다.

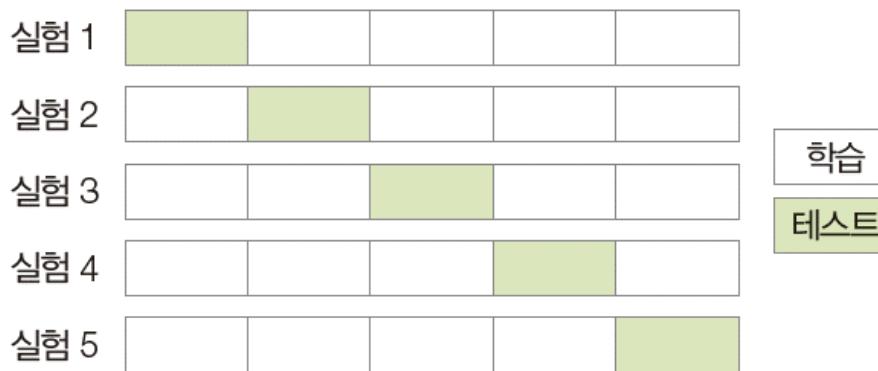


- ❖ Limitations of training/test set division
  - Likelihood of accidental high or accidental low accuracy
- ❖ k-fold cross validation
  - Use the training set divided into  $k$  subsets
  - Measure performance by learning with  $k-1$  leaving one and then leaving it
  - Increase reliability by averaging  $k$  performance

# Cross-Validation

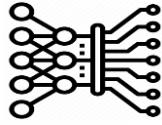


(a) 모델 선택 포함



(b) 모델 선택 제외

# Cross-Validation



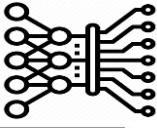
프로그램 3-6

필기 숫자 인식 – 교차 검증으로 성능 측정

```
01 from sklearn import datasets  
02 from sklearn import svm  
03 from sklearn.model_selection import cross_val_score  
04 import numpy as np  
05  
06 digit=datasets.load_digits()  
07 s=svm.SVC(gamma=0.001)  
08 accuracies=cross_val_score(s,digit.data,digit.target,cv=5) # 5-겹 교차 검증  
09  
10 print(accuracies)  
11 print("정확률(평균)=%.3f, 표준편차=%.3f"%(accuracies.mean()*100,accuracies.std()))
```

[0.97527473 0.95027624 0.98328691 0.99159664 0.95774648]

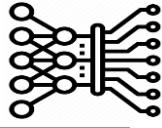
정확률(평균)=97.164, 표준편차=0.015



## ❖ Need for a variety of data

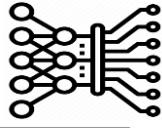
- Computers can only recognize numbers 0 and 1 unlike humans
- In other words, the object in the image should be expressed as a number, and a separate processing (weight of machine learning) should be performed to recognize only the car among the objects
  - Through this complex process, a lot of data is needed to accurately recognize various types of cars



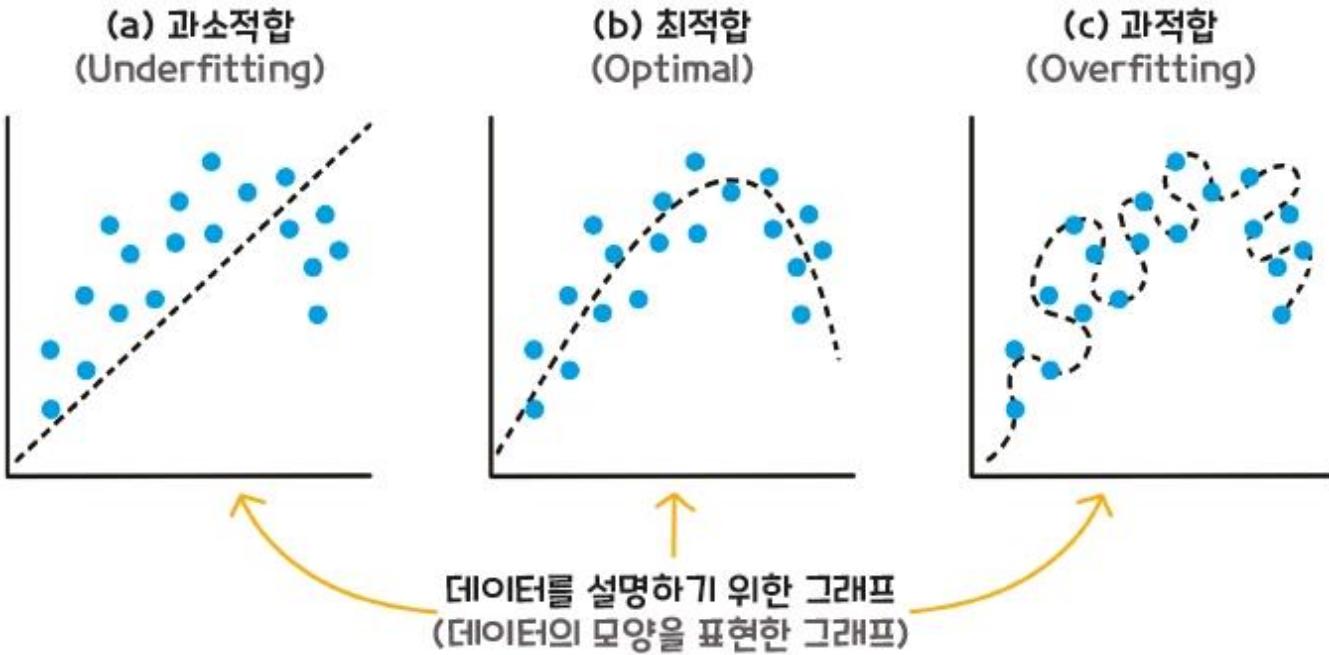


## ❖ Overfitting

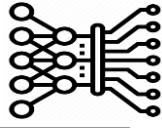
- Overlearning training data means poor performance when analyzing real data
- This occurs when data is significantly lacking compared to the complexity of the problem
  - i.e., Learning data does not cover the entire space where the problem is defined and focuses only in some cases



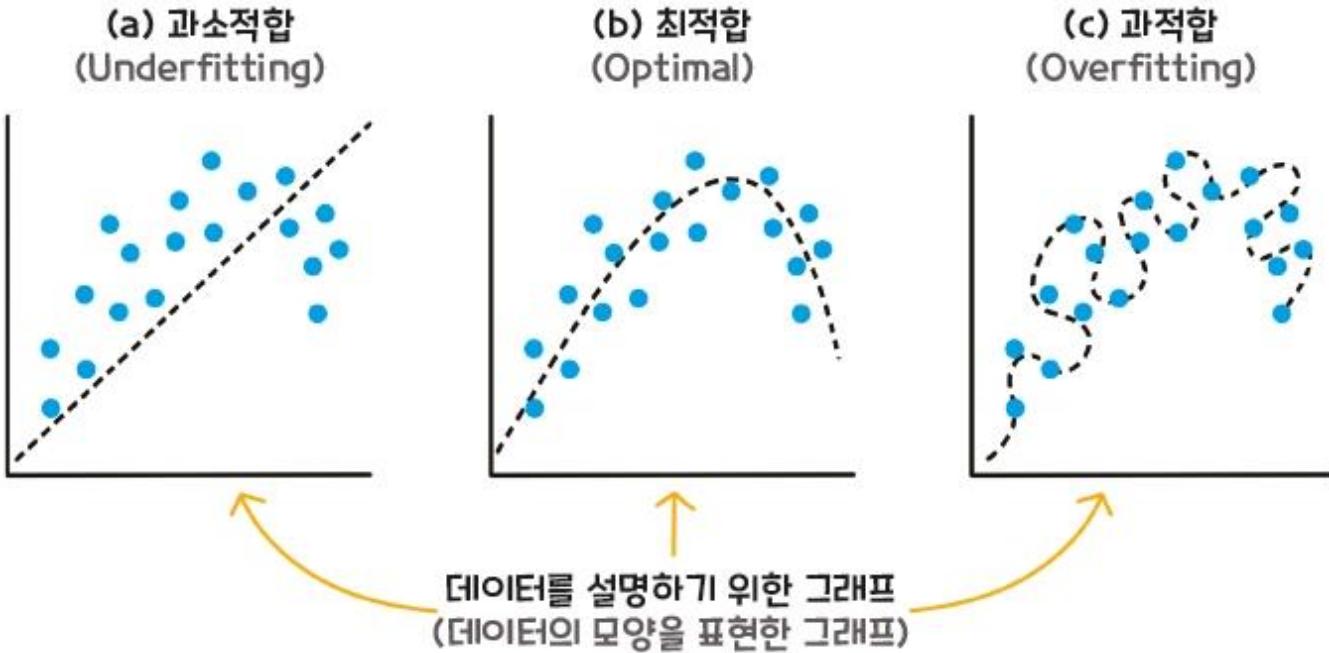
## ❖ Overfitting



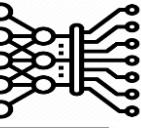
- (a) 그래프 : 변수가 0인 쪽의 데이터 몇 개는 비교적 잘 근사하지만, 일정 시점 이후 데이터는 우하향하고 있음
- (b) 그래프 : 데이터와 비슷하게 우상향하고 있어 제대로 반영하고 있다고 볼 수 있음



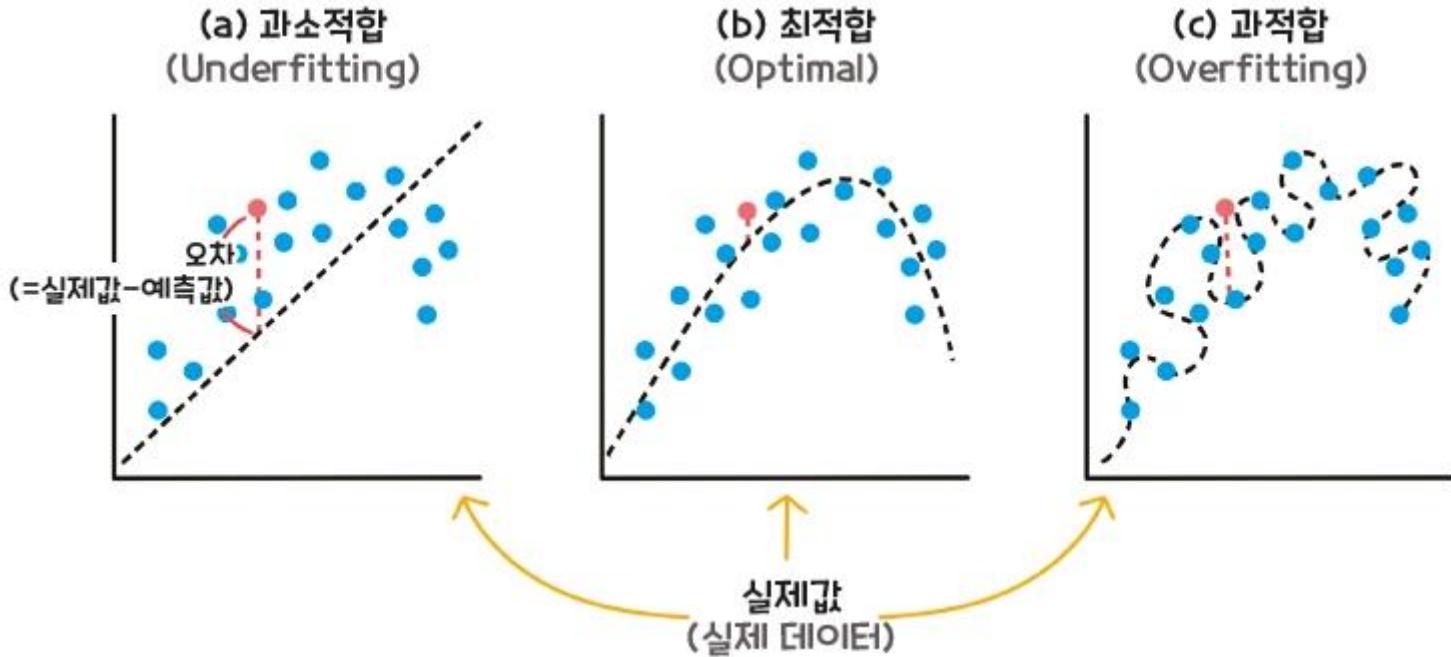
## ❖ Overfitting



- (c) 그래프 : 학습 데이터와 생성된 모델의 오차를 구해보면 0에 가까울 것임. 즉, 그래프가 모든 점을 지나고 있음
- 그렇다면 (b) 그래프보다 (c) 그래프가 더 좋은 그래프일까? 어떤 그래프가 좋은 그래프인지 알아보기 위해서는 실제 데이터값을 불러오면 됨

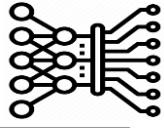


## ❖ Overfitting

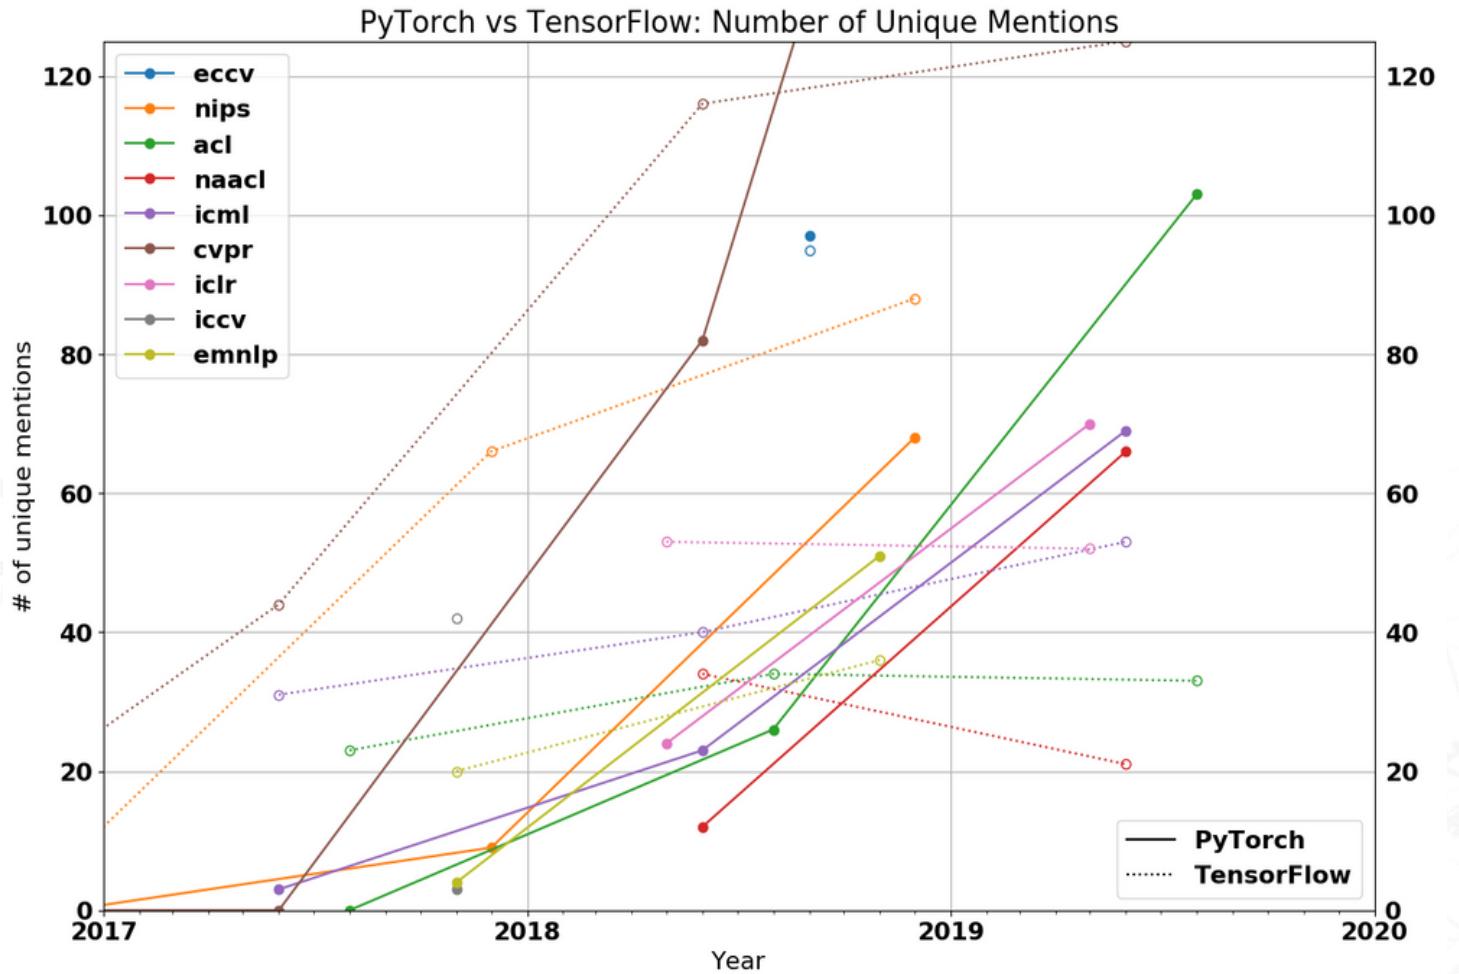


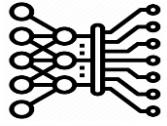
- 하나의 실제값을 불러왔을 때, 실제값과 모델이 내놓은 예측값의 차이인 오차가 가장 작은 그래프가 좋은 그래프라고 할 수 있음
- 확인해 보면 (b) 그래프의 오차가 가장 작으므로 최적합(Optimal), 즉 가장 좋은 그래프라고 할 수 있음

# Comparison of Deep Learning Libraries



## PyTorch vs. Tensorflow





## PyTorch is a Python based scientific computing package

- ✓ A cousin of LUA-based Torch framework
- ✓ Not a simple set of wrappers to support popular language



Python Support



Easy to Use API



Fast and Feels  
Native

# P Y TORCH



Actively used

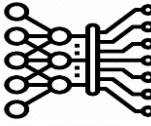


Dynamic Computation  
Graphs



Support for CUDA

# PyTorch Tutorials (1/3)



PyTorch

1.12.0+cu102

Search Tutorials

PyTorch Recipes [ + ]

Introduction to PyTorch [ - ]

Learn the Basics

- Quickstart
- Tensors
- Datasets & DataLoaders
- Transforms

Build the Neural Network

- Automatic Differentiation with `torch.autograd`
- Optimizing Model Parameters
- Save and Load the Model

Introduction to PyTorch on YouTube [ - ]

- Introduction to PyTorch - YouTube Series
- Introduction to PyTorch
- Introduction to PyTorch Tensors
- The Fundamentals of Autograd
- Building Models with PyTorch
- PyTorch TensorBoard Support
- Training with PyTorch
- Model Understanding with Captum

Learning PyTorch [ + ]

- Image and Video [ + ]
- Audio [ + ]
- Text [ + ]
- Reinforcement Learning [ + ]
- Deploying PyTorch Models in Production [ + ]
- Code Transforms with FX [ + ]
- Frontend APIs [ + ]
- Extending PyTorch [ + ]
- Model Optimization [ + ]
- Parallel and Distributed Training [ + ]
- Mobile [ + ]
- Recommendation Systems [ - ]

Tutorials > Welcome to PyTorch Tutorials

## WELCOME TO PYTORCH TUTORIALS

What's new in PyTorch tutorials?

- [Introduction to TorchRec](#)
- [Getting Started with Fully Sharded Data Parallel \(FSDP\)](#)
- [Grokking PyTorch Intel CPU Performance from First Principles](#)
- [Customize Process Group Backends Using Cpp Extensions](#)
- [Forward-mode Automatic Differentiation \(added functorch API capabilities\)](#)
- [Real Time Inference on Raspberry Pi 4 \(30 fps!\)](#)

**Learn the Basics**

Familiarize yourself with PyTorch concepts and modules. Learn how to load data, build deep neural networks, train and save your models in this quickstart guide.

[Get started with PyTorch >](#)

**PyTorch Recipes**

Bite-size, ready-to-deploy PyTorch code examples.

[Explore Recipes >](#)

All    Audio    Best Practice    C++    CUDA    Extending PyTorch    FX  
Frontend APIs    Getting Started    Image/Video    Interpretability    Memory Format  
Mobile    Model Optimization    Parallel and-Distributed-Training    Production  
Profiling    Quantization    Recommender    Reinforcement Learning    TensorBoard  
Text    TorchRec    TorchScript

Learn the Basics

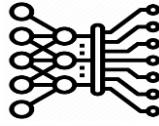
A step-by-step guide to building a complete ML workflow with PyTorch.

Getting Started

Introduction to PyTorch on YouTube

An introduction to building a complete ML workflow with PyTorch. Follows the PyTorch Beginner Series on YouTube.

<https://pytorch.org/tutorials/>



Learning PyTorch [-]

- Deep Learning with PyTorch: A 60 Minute Blitz
- Learning PyTorch with Examples
- What is `torch.nn` really?
- Visualizing Models, Data, and Training with TensorBoard

Image and Video [-]

- TorchVision Object Detection Finetuning Tutorial
- Transfer Learning for Computer Vision Tutorial
- Adversarial Example Generation
- DCGAN Tutorial
- Spatial Transformer Networks Tutorial
- Optimizing Vision Transformer Model for Deployment

Audio [-]

- Audio I/O
- Audio Resampling
- Audio Data Augmentation
- Audio Feature Extractions
- Audio Feature Augmentation
- Audio Datasets
- Speech Recognition with Wav2Vec2
- Speech Command Classification with torchaudio
- Text-to-speech with torchaudio
- Forced Alignment with Wav2Vec2

Text [-]

- Language Modeling with nn.Transformer and TorchText
- NLP From Scratch: Classifying Names with a Character-Level RNN
- NLP From Scratch: Generating Names with a Character-Level RNN
- NLP From Scratch: Translation with a Sequence to Sequence Network and Attention
- Text classification with the torchtext library
- Language Translation with nn.Transformer and torchtext

 PyTorch

1.12.0+cu102

 Search Tutorials

PyTorch Recipes [+]

Introduction to PyTorch [-]

Learn the Basics

- Quickstart
- Tensors
- Datasets & DataLoaders
- Transforms
- Build the Neural Network
- Automatic Differentiation with `torch.autograd`
- Optimizing Model Parameters
- Save and Load the Model

Introduction to PyTorch on YouTube [-]

Introduction to PyTorch - YouTube Series

Introduction to PyTorch

Introduction to PyTorch Tensors

The Fundamentals of Autograd

Building Models with PyTorch

PyTorch TensorBoard Support

Training with PyTorch

Model Understanding with Captum

Learning PyTorch [+]

Image and Video [+]

Audio [+]

Text [+]

Reinforcement Learning [+]

Deploying PyTorch Models in Production [+]

Code Transforms with FX [+]

Frontend APIs [+]

Extending PyTorch [+]

Model Optimization [+]

Parallel and Distributed Training [+]

Mobile [+]

Recommendation Systems [-]

Get Started Ecosystem Mobile Blog Tutorials Docs Resources GitHub

## WELCOME TO PYTORCH TUTORIALS

What's new in PyTorch tutorials?

- [Introduction to TorchRec](#) [-]
- [Getting Started with Fully Sharded Data Parallel \(FSDP\)](#)
- [Grokking PyTorch Intel CPU Performance from First Principles](#)
- [Customize Process Group Backends Using Cpp Extensions](#)
- [Forward-mode Automatic Differentiation \(added `functorch` API capabilities\)](#)
- [Real Time Inference on Raspberry Pi 4 \(30 fps!\)](#)

### Learn the Basics

Familiarize yourself with PyTorch concepts and modules. Learn how to load data, build deep neural networks, train and save your models in this quickstart guide.

[Get started with PyTorch >](#)

### PyTorch Recipes

Bite-size, ready-to-deploy PyTorch code examples.

[Explore Recipes >](#)



### Learn the Basics

A step-by-step guide to building a complete ML workflow with PyTorch.

[Getting Started](#)

### Introduction to PyTorch on YouTube

An introduction to building a complete ML workflow with PyTorch. Follows the PyTorch Beginner Series on YouTube.



 PyTorch

### Deploying PyTorch Models in Production [-]

Deploying PyTorch in Python via a REST API with Flask

Introduction to TorchScript

Loading a TorchScript Model in C++  
(optional) Exporting a Model from PyTorch to ONNX and Running it using ONNX Runtime

Real Time Inference on Raspberry Pi 4 (30 fps!)

### Code Transforms with FX [-]

(beta) Building a Convolution/Batch Norm fuser in FX  
(beta) Building a Simple CPU Performance Profiler with FX

### Frontend APIs [+]

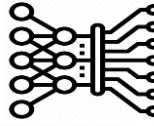
#### Extending PyTorch [-]

Double Backward with Custom Functions  
Fusing Convolution and Batch Norm using Custom Function  
Custom C++ and CUDA Extensions  
Extending TorchScript with Custom C++ Operators  
Extending TorchScript with Custom C++ Classes  
Registering a Dispatched Operator in C++  
Extending dispatcher for a new backend in C++

### Model Optimization [-]

Profiling your PyTorch Module  
PyTorch Profiler With TensorBoard  
Hyperparameter tuning with Ray Tune  
Optimizing Vision Transformer Model for Deployment  
Parametrizations Tutorial  
Pruning Tutorial  
(beta) Dynamic Quantization on an LSTM Word Language Model  
(beta) Dynamic Quantization on BERT  
(beta) Quantized Transfer Learning for Computer Vision Tutorial  
(beta) Static Quantization with Eager Mode in PyTorch  
Grokking PyTorch Intel CPU performance from first principles

<https://pytorch.org/tutorials/>



## INSTALL PYTORCH

Select your preferences and run the install command. Stable represents the most currently tested and supported version of PyTorch. This should be suitable for many users. Preview is available if you want the latest, not fully tested and supported, 1.12 builds that are generated nightly. Please ensure that you have **met the prerequisites below (e.g., numpy)**, depending on your package manager. Anaconda is our recommended package manager since it installs all dependencies. You can also [install previous versions of PyTorch](#). Note that LibTorch is only available for C++.

Additional support or warranty for some PyTorch Stable and LTS binaries are available through the [PyTorch Enterprise Support Program](#).

PyTorch Build	Stable (1.12.0)	Preview (Nightly)	LTS (1.8.2)	
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python		C++ / Java	
Compute Platform	CUDA 10.2	CUDA 11.3	CUDA 11.6	ROCM 5.1.1
Run this Command:	CUDA-10.2 PyTorch builds are no longer available for Windows, please use CUDA-11.6			

## Various PyTorch examples can use to learn and experiment

**PyTorch Examples**

Run Examples passing

<https://pytorch.org/examples/>

`pytorch/examples` is a repository showcasing examples of using PyTorch. The goal is to have curated, short, few/no dependencies *high quality* examples that are substantially different from each other that can be emulated in your existing work.

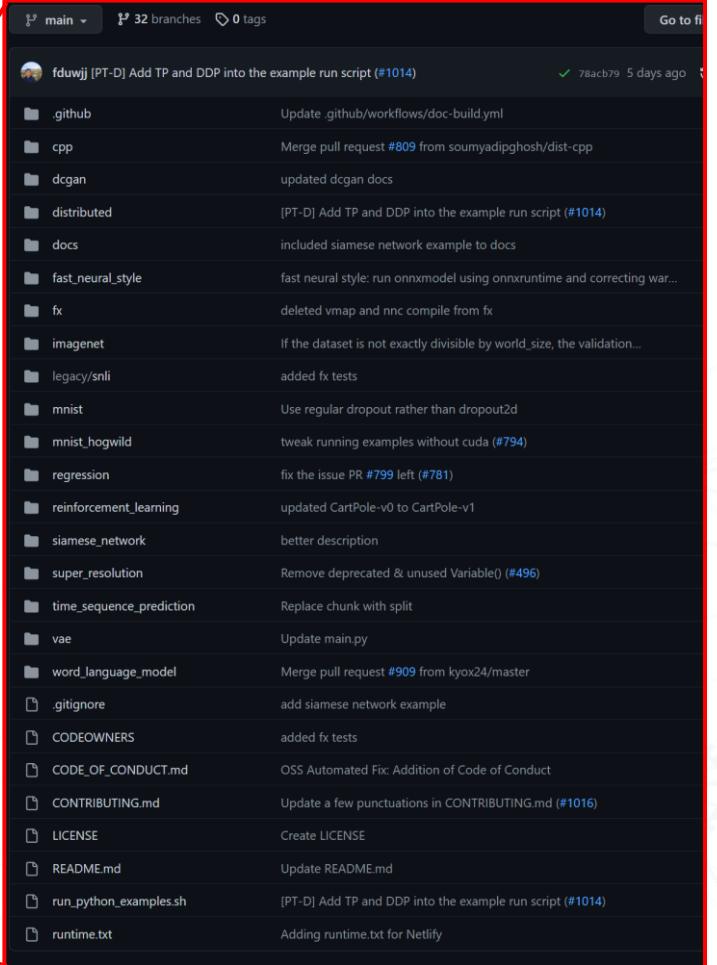
- For tutorials: <https://github.com/pytorch/tutorials>
- For changes to pytorch.org: <https://github.com/pytorch/pytorch.github.io>
- For a general model hub: <https://pytorch.org/hub/> or <https://huggingface.co/models>
- For recipes on how to run PyTorch in production: <https://github.com/facebookresearch/recipes>
- For general Q&A and support: <https://discuss.pytorch.org/>

**Available models**

- Image classification (MNIST) using Convnets
- Word-level Language Modeling using RNN and Transformer
- Training Imagenet Classifiers with Popular Networks
- Generative Adversarial Networks (DCGAN)
- Variational Auto-Encoders
- Superresolution using an efficient sub-pixel convolutional neural network
- Hogwild training of shared ConvNets across multiple processes on MNIST
- Training a CartPole to balance in OpenAI Gym with actor-critic
- Natural Language Inference (SNLI) with GloVe vectors, LSTMs, and torchtext
- Time sequence prediction - use an LSTM to learn Sine waves
- Implement the Neural Style Transfer algorithm on images
- Reinforcement Learning with Actor Critic and REINFORCE algorithms on OpenAI gym
- PyTorch Module Transformations using fx
- Distributed PyTorch examples with Distributed Data Parallel and RPC
- Several examples illustrating the C++ Frontend

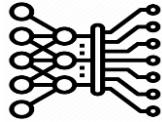
Additionally, a list of good examples hosted in their own repositories:

- Neural Machine Translation using sequence-to-sequence RNN with attention (OpenNMT)



The screenshot shows the GitHub repository for PyTorch Examples. The main branch is 'main' with 32 branches and 0 tags. A recent commit by fduwjj was pushed 5 days ago, adding TP and DDP into the example run script (#1014). The repository contains several sub-directories and files, each with a brief description of its purpose. Some commits are linked to specific pull requests, such as #809 and #799.

File / Directory	Description
.github	Update .github/workflows/doc-build.yml
cpp	Merge pull request #809 from soumyadipgosh/dist-cpp
dcgan	updated dcgan docs
distributed	[PT-D] Add TP and DDP into the example run script (#1014)
docs	included siamese network example to docs
fast_neural_style	fast neural style: run onnxmodel using onnxruntime and correcting war...
fx	deleted vmap and rnc compile from fx
imagenet	If the dataset is not exactly divisible by world_size, the validation...
legacy/snli	added fx tests
mnist	Use regular dropout rather than dropout2d
mnist_hogwild	tweak running examples without cuda (#794)
regression	fix the issue PR #799 left (#781)
reinforcement_learning	updated CartPole-v0 to CartPole-v1
siamese_network	better description
super_resolution	Remove deprecated & unused Variable() (#496)
time_sequence_prediction	Replace chunk with split
vae	Update main.py
word_language_model	Merge pull request #909 from kyox24/master
.gitignore	add siamese network example
CODEOWNERS	added fx tests
CODE_OF_CONDUCT.md	OSS Automated Fix: Addition of Code of Conduct
CONTRIBUTING.md	Update a few punctuations in CONTRIBUTING.md (#1016)
LICENSE	Create LICENSE
README.md	Update README.md
run_python_examples.sh	[PT-D] Add TP and DDP into the example run script (#1014)
runtime.txt	Adding runtime.txt for Netlify



## Various PyTorch examples can use to learn and experiment

### Image Classification Using ConvNets

This example demonstrates how to run image classification with [Convolutional Neural Networks ConvNets](#) on the [MNIST](#) database.

[GO TO EXAMPLE](#) ↗

### Measuring Similarity using Siamese Network

This example demonstrates how to measure similarity between two images using [Siamese network](#) on the [MNIST](#) database.

[GO TO EXAMPLE](#) ↗

### Super-resolution Using an Efficient Sub-Pixel CNN

This example demonstrates how to use the sub-pixel convolution layer described in [Real-Time Single Image and Video Super-Resolution Using an Efficient Sub-Pixel Convolutional Neural Network](#) paper. This example trains a super-resolution network on the [BSD300](#) dataset.

[GO TO EXAMPLE](#) ↗

### HOGWILD! Training of Shared ConvNets

[HOGWILD!](#) is a scheme that allows Stochastic Gradient Descent (SGD) parallelization without memory locking. This example demonstrates how to perform HOGWILD! training of shared ConvNets on MNIST.

[GO TO EXAMPLE](#) ↗

### Word-level Language Modeling using RNN and Transformer

This example demonstrates how to train a multi-layer [recurrent neural network \(RNN\)](#), such as Elman, GRU, or LSTM, or Transformer on a language modeling task by using the [Wikitext-2](#) dataset.

[GO TO EXAMPLE](#) ↗

### Training ImageNet Classifiers

This example demonstrates how you can train some of the most popular model architectures, including [ResNet](#), [AlexNet](#), and [VGG](#) on the [ImageNet](#) dataset.

[GO TO EXAMPLE](#) ↗

### Training a CartPole to balance in OpenAI Gym with actor-critic

This reinforcement learning tutorial demonstrates how to train a CartPole to balance in the [OpenAI Gym](#) toolkit by using the [Actor-Critic](#) method.

[GO TO EXAMPLE](#) ↗

### Time Sequence Prediction

This beginner example demonstrates how to use [LSTMCell](#) to learn sine wave signals to predict the signal values in the future.

[GO TO EXAMPLE](#) ↗

### Generative Adversarial Networks (DCGAN)

This example implements the [Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks](#) paper.

[GO TO EXAMPLE](#) ↗

### Variational Auto-Encoders

This example implements the [Auto-Encoding Variational Bayes](#) paper with [ReLUs](#) and the Adam optimizer.

[GO TO EXAMPLE](#) ↗

### Implement the Neural Style Transfer algorithm on images

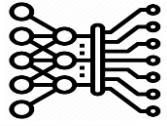
This tutorial demonstrates how you can use PyTorch's implementation of the [Neural Style Transfer \(NST\)](#) algorithm on images.

[GO TO EXAMPLE](#) ↗

### PyTorch Module Transformations using fx

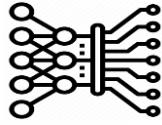
This set of examples demonstrates the [torch.fx](#) toolkit. For more information about [torch.fx](#), see [torch.fx Overview](#).

[GO TO EXAMPLE](#) ↗



## Various PyTorch examples can use to learn and experiment

<b>Image Classification Using ConvNets</b>  This example demonstrates how to run image classification with <a href="#">Convolutional Neural Networks ConvNets</a> on the <a href="#">MNIST</a> database.  <a href="#">GO TO EXAMPLE</a> ↗	<b>Measuring Similarity using Siamese Network</b>  This example demonstrates how to measure similarity between two images using Siamese network on the MNIST database.  <a href="#">GO TO EXAMPLE</a> ↗	<b>Super-resolution Using an Efficient Sub-Pixel CNN</b>  This example demonstrates how to use the sub-pixel convolution layer described in <a href="#">Real-Time Single Image and Video Super-Resolution Using an Efficient Sub-Pixel Convolutional Neural Network</a> paper. This example trains a super-resolution network on the BSD300 dataset.  <a href="#">GO TO EXAMPLE</a> ↗	<b>HOGWILD! Training of Shared ConvNets</b>  HOGWILD! is a scheme that allows Stochastic Gradient Descent (SGD) parallelization without memory locking. This example demonstrates how to perform HOGWILD! training of shared ConvNets on MNIST.  <a href="#">GO TO EXAMPLE</a> ↗
<b>Word-level Language Modeling using RNN and Transformer</b>  This example demonstrates how to train a multi-layer recurrent neural network (RNN), such as Elman, GRU, or LSTM, or Transformer on a language modeling task by using the Wikitext-2 dataset.  <a href="#">GO TO EXAMPLE</a> ↗	<b>Training ImageNet Classifiers</b>  This example demonstrates how you can train some of the most popular model architectures, including ResNet, AlexNet, and VGG on the ImageNet dataset.  <a href="#">GO TO EXAMPLE</a> ↗	<b>Training a CartPole to balance in OpenAI Gym with actor-critic</b>  This reinforcement learning tutorial demonstrates how to train a CartPole to balance in the OpenAI Gym toolkit by using the Actor-Critic method.  <a href="#">GO TO EXAMPLE</a> ↗	<b>Time Sequence Prediction</b>  This beginner example demonstrates how to use LSTMCell to learn sine wave signals to predict the signal values in the future.  <a href="#">GO TO EXAMPLE</a> ↗
<b>Generative Adversarial Networks (DCGAN)</b>  This example implements the Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks paper.  <a href="#">GO TO EXAMPLE</a> ↗	<b>Variational Auto-Encoders</b>  This example implements the Auto-Encoding Variational Bayes paper with ReLUs and the Adam optimizer.  <a href="#">GO TO EXAMPLE</a> ↗	<b>Implement the Neural Style Transfer algorithm on images</b>  This tutorial demonstrates how you can use PyTorch's implementation of the Neural Style Transfer (NST) algorithm on images.  <a href="#">GO TO EXAMPLE</a> ↗	<b>PyTorch Module Transformations using fx</b>  This set of examples demonstrates the torch.fx toolkit. For more information about <code>torch.fx</code> , see <a href="#">torch.fx Overview</a> .  <a href="#">GO TO EXAMPLE</a> ↗



## Implement of CNN architecture for visual recognition

CS231n Convolutional Neural Networks for Visual Recognition

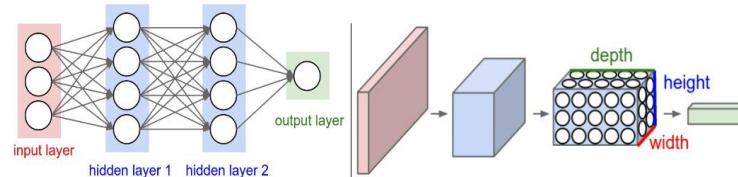
Table of Contents:

- Architecture Overview
- ConvNet Layers
  - Convolutional Layer
  - Pooling Layer
  - Normalization Layer
  - Fully-Connected Layer
  - Converting Fully-Connected Layers to Convolutional Layers
- ConvNet Architectures
  - Layer Patterns
  - Layer Sizing Patterns
  - Case Studies (LeNet / AlexNet / ZFNet / GoogLeNet / VGGNet)
  - Computational Considerations
- Additional References

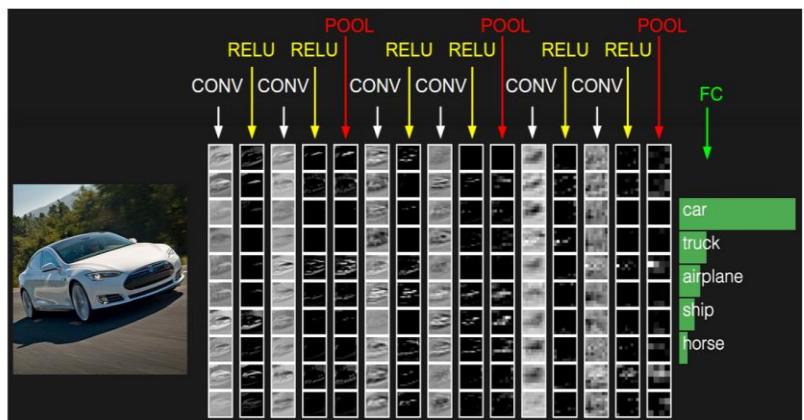
### Convolutional Neural Networks (CNNs / ConvNets)

Convolutional Neural Networks are very similar to ordinary Neural Networks from the previous chapter: they are made up of neurons that have learnable weights and biases. Each neuron receives some inputs, performs a dot product and optionally follows it with a non-linearity. The whole network still expresses a single differentiable score function: from the raw image pixels on one end to class scores at the other. And they still have a loss function (e.g. SVM/Softmax) on the last (fully-connected) layer and all the tips/tricks we developed for learning regular Neural Networks still apply.

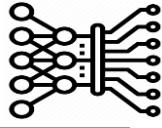
So what changes? ConvNet architectures make the explicit assumption that the inputs are images, which allows us to encode certain properties into the architecture. These then make the forward function more efficient to implement and vastly reduce the amount of parameters in the network.



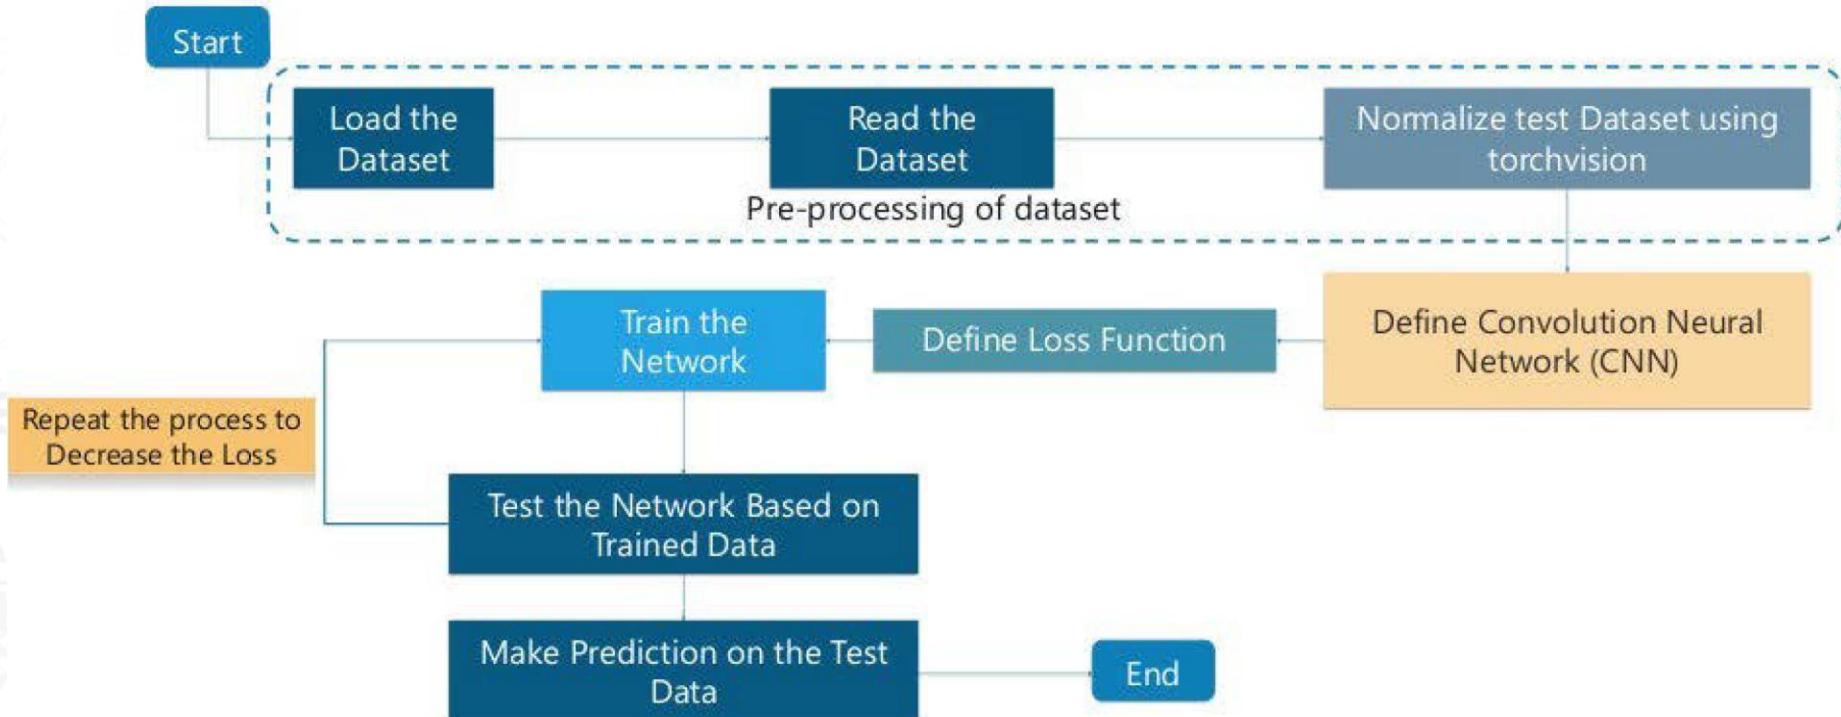
Left: A regular 3-layer Neural Network. Right: A ConvNet arranges its neurons in three dimensions (width, height, depth), as visualized in one of the layers. Every layer of a ConvNet transforms the 3D input volume to a 3D output volume of neuron activations. In this example, the red input layer holds the image, so its width and height would be the dimensions of the image, and the depth would be 3 (Red, Green, Blue channels).

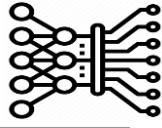


The activations of an example ConvNet architecture. The initial volume stores the raw image pixels (left) and the last volume stores the class scores (right). Each volume of activations along the processing path is shown as a column. Since it's difficult to visualize 3D volumes, we lay out each volume's slices in rows. The last layer volume holds the scores for each class, but here we only visualize the sorted top 5 scores, and print the labels of each one. The full [web-based demo](#) is shown in the header of our website. The architecture shown here is a tiny VGG Net, which we will discuss later.



## Developing flow diagram





## Step 1: Loading and normalizing MNIST image dataset

Load and normalize dataset

Define CNN

Define loss function

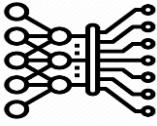
Train the network

Validate the network

Update the weights

```
1  from __future__ import print_function
2  import argparse
3  import torch
4  import torch.nn as nn
5  import torch.nn.functional as F
6  import torch.optim as optim
7  from torchvision import datasets, transforms
8  from torch.optim.lr_scheduler import StepLR
9
```

```
116      dataset1 = datasets.MNIST('../data', train=True, download=True,
117                                  transform=transform)
118      dataset2 = datasets.MNIST('../data', train=False,
119                                  transform=transform)
```



## Step 2: Define a convolution neural network

Load and normalize dataset

Define CNN

Define loss function

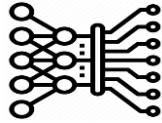
Train the network

Validate the network

Update the weights

```
11  class Net(nn.Module):
12      def __init__(self):
13          super(Net, self).__init__()
14          self.conv1 = nn.Conv2d(1, 32, 3, 1)
15          self.conv2 = nn.Conv2d(32, 64, 3, 1)
16          self.dropout1 = nn.Dropout(0.25)
17          self.dropout2 = nn.Dropout(0.5)
18          self.fc1 = nn.Linear(9216, 128)
19          self.fc2 = nn.Linear(128, 10)
20
21      def forward(self, x):
22          x = self.conv1(x)
23          x = F.relu(x)
24          x = self.conv2(x)
25          x = F.relu(x)
26          x = F.max_pool2d(x, 2)
27          x = self.dropout1(x)
28          x = torch.flatten(x, 1)
29          x = self.fc1(x)
30          x = F.relu(x)
31          x = self.dropout2(x)
32          x = self.fc2(x)
33          output = F.log_softmax(x, dim=1)
34
35      return output
```

<https://github.com/pytorch/examples/blob/main/mnist/main.py>



## Step 3: Define a loss function and optimizer

Load and normalize dataset

Define CNN

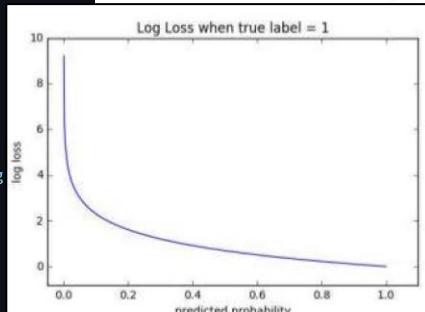
Define loss function

Train the network

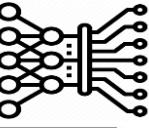
Validate the network

Update the weights

```
73 def main():
74     # Training settings
75     parser = argparse.ArgumentParser(description='PyTorch MNIST Example')
76     parser.add_argument('--batch-size', type=int, default=64, metavar='N',
77                         help='input batch size for training (default: 64)')
78     parser.add_argument('--test-batch-size', type=int, default=1000, metavar='N',
79                         help='input batch size for testing (default: 1000)')
80     parser.add_argument('--epochs', type=int, default=14, metavar='N',
81                         help='number of epochs to train (default: 14)')
82     parser.add_argument('--lr', type=float, default=1.0, metavar='LR',
83                         help='learning rate (default: 1.0)')
84     parser.add_argument('--gamma', type=float, default=0.7, metavar='M',
85                         help='Learning rate step gamma (default: 0.7)')
86     parser.add_argument('--no-cuda', action='store_true', default=False,
87                         help='disables CUDA training')
88     parser.add_argument('--dry-run', action='store_true', default=False,
89                         help='quickly check a single pass')
90     parser.add_argument('--seed', type=int, default=1, metavar='S',
91                         help='random seed (default: 1)')
92     parser.add_argument('--log-interval', type=int, default=10, metavar='N',
93                         help='how many batches to wait before logging training')
94     parser.add_argument('--save-model', action='store_true', default=False,
95                         help='For saving the current Model')
96     args = parser.parse_args()
97     use_cuda = not args.no_cuda and torch.cuda.is_available()
98
99     torch.manual_seed(args.seed)
100
101    device = torch.device("cuda" if use_cuda else "cpu")
```



Classification Cross-Entropy Loss



## Step 4: Train the network

Load and normalize dataset

Define CNN

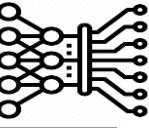
Define loss function

Train the network

Validate the network

Update the weights

```
37  def train(args, model, device, train_loader, optimizer, epoch):
38      model.train()
39      for batch_idx, (data, target) in enumerate(train_loader):
40          data, target = data.to(device), target.to(device)
41          optimizer.zero_grad()
42          output = model(data)
43          loss = F.nll_loss(output, target)
44          loss.backward()
45          optimizer.step()
46          if batch_idx % args.log_interval == 0:
47              print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
48                  epoch, batch_idx * len(data), len(train_loader.dataset),
49                  100. * batch_idx / len(train_loader), loss.item()))
50          if args.dry_run:
51              break
```



## Step 5: Validate the network

Load and normalize dataset

Define CNN

Define loss function

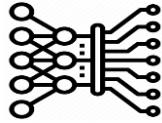
Train the network

Validate the network

Update the weights

```
53
54     def test(model, device, test_loader):
55         model.eval()
56         test_loss = 0
57         correct = 0
58         with torch.no_grad():
59             for data, target in test_loader:
60                 data, target = data.to(device), target.to(device)
61                 output = model(data)
62                 test_loss += F.nll_loss(output, target, reduction='sum').item() # sum up batch loss
63                 pred = output.argmax(dim=1, keepdim=True) # get the index of the max log-probability
64                 correct += pred.eq(target.view_as(pred)).sum().item()
65
66         test_loss /= len(test_loader.dataset)
67
68         print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)\n'.format(
69             test_loss, correct, len(test_loader.dataset),
70             100. * correct / len(test_loader.dataset)))
```

# Image Classification on the MNIST (8/8)



## Step 5: Update the weights and parameter tuning and save the model

Load and normalize dataset

Define CNN

Define loss function

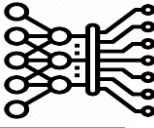
Train the network

Validate the network

Update the weights

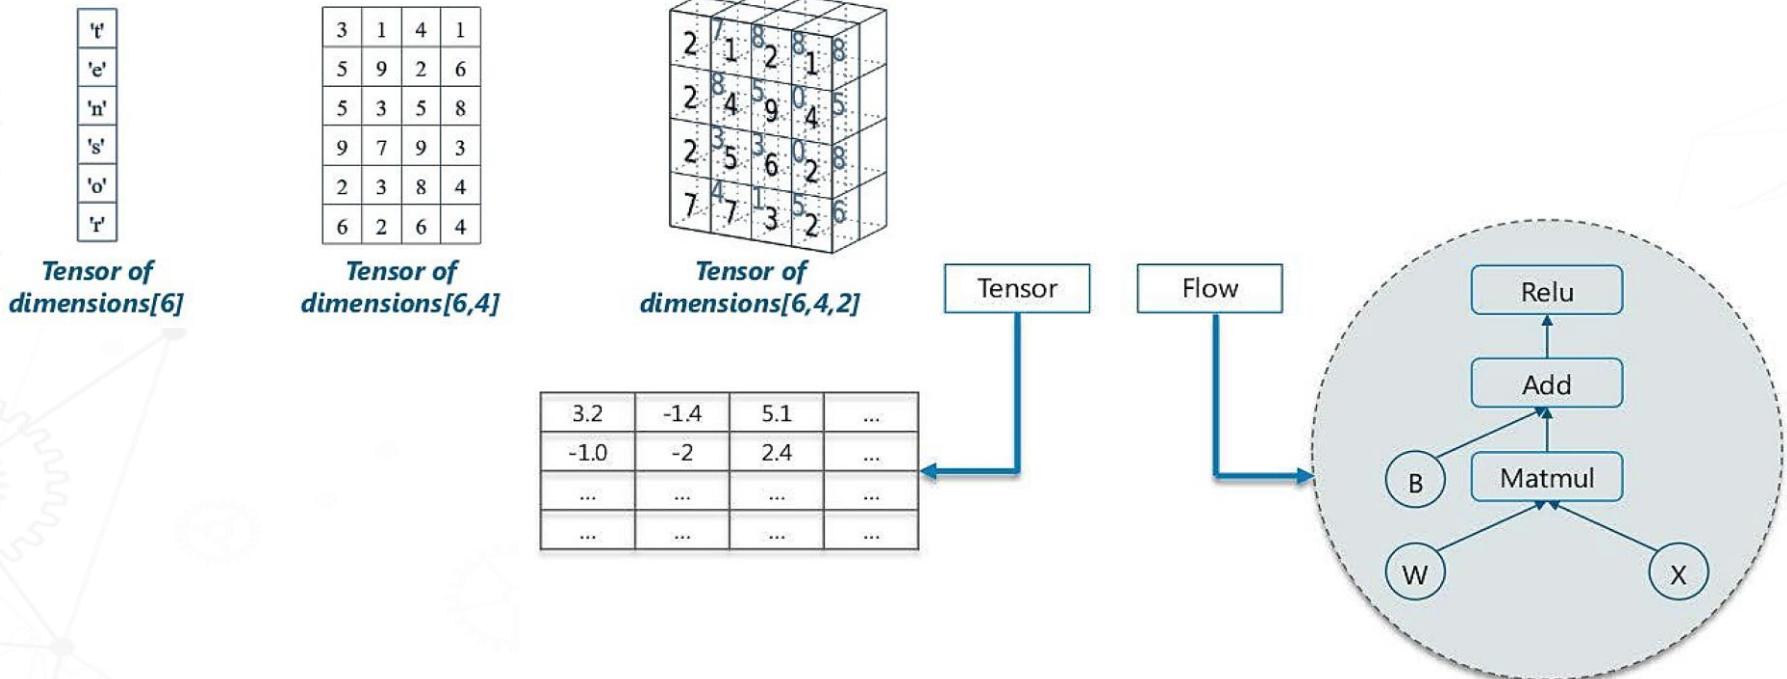
```
120     train_loader = torch.utils.data.DataLoader(dataset1, **train_kwargs)
121     test_loader = torch.utils.data.DataLoader(dataset2, **test_kwargs)
122
123     model = Net().to(device)
124     optimizer = optim.Adadelta(model.parameters(), lr=args.lr)
125
126     scheduler = StepLR(optimizer, step_size=1, gamma=args.gamma)
127     for epoch in range(1, args.epochs + 1):
128         train(args, model, device, train_loader, optimizer, epoch)
129         test(model, device, test_loader)
130         scheduler.step()
131
132     if args.save_model:
133         torch.save(model.state_dict(), "mnist_cnn.pt")
134
135
136 if __name__ == '__main__':
137     main()
```

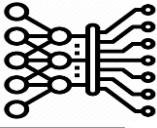
<https://github.com/pytorch/examples/blob/main/mnist/main.py>



## The standard way of representing data in deep learning

- ✓ Just multidimensional arrays, an extension of two-dimensional matrices to data with higher dimension
- ✓ In Tensorflow, computation is approached as a dataflow graph

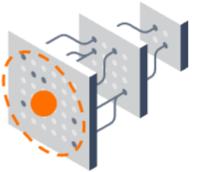




## Why TensorFlow

TensorFlow is an end-to-end open source platform for machine learning. It has a comprehensive, flexible ecosystem of tools, libraries and community resources that lets researchers push the state-of-the-art in ML and developers easily build and deploy ML powered applications.

About →



### Easy model building

Build and train ML models easily using intuitive high-level APIs like Keras with eager execution, which makes for immediate model iteration and easy debugging.



### Robust ML production anywhere

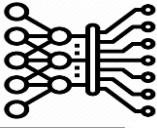
Easily train and deploy models in the cloud, on-prem, in the browser, or on-device no matter what language you use.



### Powerful experimentation for research

A simple and flexible architecture to take new ideas from concept to code, to state-of-the-art models, and to publication faster.

# TensorFlow Tutorials (1/9)



The screenshot shows the TensorFlow website's "Install" section. The left sidebar has a "Install TensorFlow" button highlighted. The main content area is titled "Install TensorFlow 2". It lists supported systems and provides links for "Download a package" (using pip or Docker), "Run a TensorFlow container" (using Docker), "Google Colab" (a Jupyter notebook environment), and "Build your first ML app" (for web developers using TensorFlow.js or mobile developers using TensorFlow Lite). A search bar is at the top right.

Install

Install TensorFlow

Packages  
pip  
Docker

Additional setup  
GPU device plugins  
Problems

Build from source  
Linux / macOS  
Windows  
SIG Build

Language bindings  
Java  
Java (legacy)  
C  
Go

Install TensorFlow 2

TensorFlow is tested and supported on the following 64-bit systems:

- Python 3.7–3.10
- Ubuntu 16.04 or later
- Windows 7 or later (with C++ redistributable)
- macOS 10.12.6 (Sierra) or later (no GPU support)
- WSL2 via Windows 10 19044 or higher including GPUs (Experimental)

Download a package

Install TensorFlow with Python's `pip` package manager.

★ TensorFlow 2 packages require a `pip` version >19.0 (or >20.3 for macOS).

```
# Requires the latest pip
$ pip install --upgrade pip

# Current stable release for CPU and GPU
$ pip install tensorflow

# Or try the preview build (unstable)
$ pip install tf-nightly
```

Official packages available for Ubuntu, Windows, and macOS.

Read the `pip` install guide

Run a TensorFlow container

The TensorFlow Docker images are already configured to run TensorFlow. A Docker container runs in a virtual environment and is the easiest way to set up GPU support.

```
$ docker pull tensorflow/tensorflow:latest # Download latest stable image
$ docker run -it -p 8888:8888 tensorflow/tensorflow:latest-jupyter # Start Jupyter server
```

Read the Docker install guide

Google Colab: An easy way to learn and use TensorFlow

No install necessary—run the TensorFlow tutorials directly in the browser with Colaboratory, a Google research project created to help disseminate machine learning education and research. It's a Jupyter notebook environment that requires no setup to use and runs entirely in the cloud. Read the blog post.

Build your first ML app

Create and deploy TensorFlow models on web and mobile.

Web developers

TensorFlow.js is a WebGL accelerated, JavaScript library to train and deploy ML models in the browser, Node.js, mobile, and more.

Mobile developers

TensorFlow Lite is a lightweight solution for mobile and embedded devices.

<https://www.tensorflow.org/>

## GPU device plugins

★ Note: This page is for non-NVIDIA® GPU devices. For NVIDIA® GPU support, go to the [Install TensorFlow with pip guide](#).

TensorFlow's [pluggable device\(external\)](#) architecture adds new device support as separate plug-in packages that are installed alongside the official TensorFlow package.

The mechanism requires no device-specific changes in the TensorFlow code. It relies on C APIs to communicate with the TensorFlow binary in a stable manner. Plug-in developers maintain separate code repositories and distribution packages for their plugins and are responsible for testing their devices.

### Use device plugins

To use a particular device, like one would a native device in TensorFlow, users only have to install the device plugin package for that device. The following code snippet shows how the plugin for a new demonstration device, Awesome Processing Unit (APU), is installed and used. For simplicity, this sample APU plug-in only has one custom kernel for ReLU:

```
# Install the APU example plug-in package
$ pip install tensorflow-apu-0.0.1-cp36-cp36m-linux_x86_64.whl
...
Successfully installed tensorflow-apu-0.0.1
```

With the plug-in installed, test that the device is visible and run an operation on the new APU device:

```
import tensorflow as tf # TensorFlow registers PluggableDevices here.
tf.config.list_physical_devices() # APU device is visible to TensorFlow.
[PhysicalDevice(name='/physical_device:CPU:0', device_type='CPU'), PhysicalDevice(name='/physical_device:'

a = tf.random.normal(shape=[5], dtype=tf.float32) # Runs on CPU.
b = tf.nn.relu(a) # Runs on APU.

with tf.device("/APU:0"): # Users can also use 'with tf.device' syntax.
    c = tf.nn.relu(a) # Runs on APU.

with tf.device("/CPU:0"):
    c = tf.nn.relu(a) # Runs on CPU.

@tf.function # Defining a tf.function
def run():
    d = tf.random.uniform(shape=[100], dtype=tf.float32) # Runs on CPU.
    e = tf.nn.relu(d) # Runs on APU.

run() # PluggableDevices also work with tf.function and graph mode.
```



# TensorFlow Tutorials (3/9)



## TensorFlow is an end-to-end open source platform for machine learning

TensorFlow makes it easy for beginners and experts to create machine learning models. See the sections below to get started.

[See tutorials](#)[See the guide](#)

Tutorials show you how to use TensorFlow with complete, end-to-end examples.

Guides explain the concepts and components of TensorFlow.



### For beginners

The best place to start is with the user-friendly Sequential API. You can create models by plugging together building blocks. Run the "Hello World" example below, then visit the [tutorials](#) to learn more.

To learn ML, check out our [education page](#). Begin with curated curriculums to improve your skills in foundational ML areas.

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

[Run code now](#)[Try in Google's interactive notebook](#)

### For experts

The Subclassing API provides a define-by-run interface for advanced research. Create a class for your model, then write the forward pass imperatively. Easily author custom layers, activations, and training loops. Run the "Hello World" example below, then visit the [tutorials](#) to learn more.

```
class MyModel(tf.keras.Model):
    def __init__(self):
        super(MyModel, self).__init__()
        self.conv1 = Conv2D(32, 3, activation='relu')
        self.flatten = Flatten()
        self.d1 = Dense(128, activation='relu')
        self.d2 = Dense(10, activation='softmax')

    def call(self, x):
        x = self.conv1(x)
        x = self.flatten(x)
        x = self.d1(x)
        return self.d2(x)

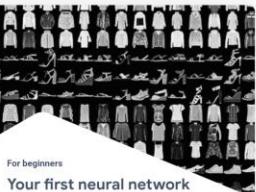
model = MyModel()

with tf.GradientTape() as tape:
    logits = model(images)
    loss_value = loss(logits, labels)
grads = tape.gradient(loss_value, model.trainable_variables)
optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

[Run code now](#)[Try in Google's interactive notebook](#)

## Solutions to common problems

Explore step-by-step tutorials to help you with your projects.



### Your first neural network

Train a neural network to classify images of clothing, like sneakers and shirts, in this fast-paced overview of a complete TensorFlow program.



### Generative adversarial networks

Train a generative adversarial network to generate images of handwritten digits, using the Keras Subclassing API.



### Neural machine translation with attention

Train a sequence-to-sequence model for Spanish to English translation using the Keras Subclassing API.

## News & announcements

Check out our [blog](#) for additional updates, and subscribe to our [TensorFlow newsletter](#) to get the latest announcements sent directly to your inbox.

[Sign up ↗](#)

June 28, 2022

### Bringing Machine Learning to every developer's toolbox

With the release of the recent Stack Overflow Developer Survey, we're delighted to see the growth of TensorFlow as the most-used...

[Read the blog →](#)

June 10, 2022

### Adding Quantization-aware Training and Pruning to the TensorFlow Model Garden

The TensorFlow model optimization toolkit (TFMOT) provides modern optimization techniques such as quantization...

[Read the blog →](#)

June 2, 2022

### New documentation on tensorflow.org

As Google I/O took place, we published a lot of exciting new docs on tensorflow.org, including updates to model parallelism and...

[Read the blog →](#)

<https://www.tensorflow.org/>

# TensorFlow Tutorials (4/9)



The screenshot shows the TensorFlow Core Tutorials page. The navigation bar includes links for Install, Learn (selected), API, Resources, Community, and Why TensorFlow, along with a search bar. The main content area is titled "TensorFlow Core" and "TensorFlow Core Tutorials". It provides instructions for running tutorials in Google Colab and lists categories for beginners and experts.

**TensorFlow tutorials**

- Quickstart for beginners
- Quickstart for experts

**BEGINNER**

- ML basics with Keras
- Load and preprocess data

**ADVANCED**

- Customization
- Distributed training
- Images
- Text
- Audio
- Structured data
- Generative
- Model Understanding
- Reinforcement learning
- tf.Estimator

The TensorFlow tutorials are written as Jupyter notebooks and run directly in Google Colab—a hosted notebook environment that requires no setup. Click the [Run in Google Colab](#) button.

## For beginners

The best place to start is with the user-friendly Keras sequential API. Build models by plugging together building blocks. After these tutorials, read the [Keras guide](#).

**Beginner quickstart**  
This "Hello, World!" notebook shows the Keras Sequential API and `model.fit`.

**Keras basics**  
This notebook collection demonstrates basic machine learning tasks using Keras.

**Load data**  
These tutorials use `tf.data` to load various data formats and build input pipelines.

## For experts

The Keras functional and subclassing APIs provide a define-by-run interface for customization and advanced research. Build your model, then write the forward and backward pass. Create custom layers, activations, and training loops.

**Advanced quickstart**  
This "Hello, World!" notebook uses the Keras subclassing API and a custom training loop.

**Customization**  
This notebook collection shows how to build custom layers and training loops in TensorFlow.

**Distributed training**  
Distribute your model training across multiple GPUs, multiple machines or TPUs.

The Advanced section has many instructive notebooks examples, including [Neural machine translation](#), [Transformers](#), and [CycleGAN](#).

<https://www.tensorflow.org/>

# TensorFlow Tutorials (5/9)



Screenshot of the TensorFlow Core Tutorials page. The URL is <https://www.tensorflow.org/tutorials>.

The page shows the "TensorFlow tutorials" section under the "TensorFlow Core" category. The "Tutorials" tab is selected. A red box highlights the "ML basics with Keras" and "Load and preprocess data" sections in the sidebar.

**BEGINNER**

**ML basics with Keras**

- Load and preprocess data
  - Images
  - CSV
  - NumPy
  - pandas.DataFrame
  - TFRecord and tf.Example
  - Additional formats with tf.io
  - Text
  - More text loading
    - Unicode
    - Subword Tokenization

The Keras functional and subclassing APIs provide a define-by-run interface for customization and advanced research. Build your model, then write the forward and backward pass. Create custom layers, activations, and training loops.

**Advanced quickstart**  
This "Hello, World!" notebook uses the Keras subclassing API and a custom training loop.

**Customization**  
This notebook collection shows how to build custom layers and training loops in TensorFlow.

**Distributed training**  
Distribute your model training across multiple GPUs, multiple machines or TPUs.

**Load data**  
These tutorials use `tf.data` to load various data formats and build input pipelines.

The Advanced section has many instructive notebooks examples, including [Neural machine translation](#), [Transformers](#), and [CycleGAN](#).

# TensorFlow Tutorials (6/9)



TensorFlow

Install Learn API Resources Community Why TensorFlow Search

TensorFlow Core

Overview Tutorials Guide Migrate to TF2

Filter

TensorFlow tutorials

- Quickstart for beginners
- Quickstart for experts

BEGINNER

- ML basics with Keras
- Load and preprocess data

ADVANCED

- Customization
- Distributed training
- Images
- Text
- Audio
- Structured data
- Generative
- Model Understanding
- Reinforcement learning
- tf.Estimator

Images

- Convolutional Neural Network
- Image classification
- Transfer learning and fine-tuning
- Transfer learning with TF Hub
- Data Augmentation
- Image segmentation
- Object detection with TF Hub

Text

- Word embeddings
- Word2Vec
- Text classification with an RNN
- Classify Text with BERT
- Solve GLUE tasks using BERT on TPU
- Neural machine translation with attention
- Image captioning

Load data

These tutorials use `tf.data` to load various data formats and build input pipelines.

Distributed training

Distribute your model training across multiple GPUs, multiple machines or TPUs.

The Advanced section has many instructive notebooks examples, including [Neural machine translation](#), [Transformers](#), and [CycleGAN](#).

# TensorFlow Tutorials (7/9)



The screenshot shows the TensorFlow Core Tutorials page. The navigation bar includes links for Install, Learn (selected), API, Resources, Community, and Why TensorFlow. A search bar is at the top right. The main content area is titled "TensorFlow Core" and contains sections for "Overview", "Tutorials" (selected), "Guide", and "Migrate to TF2". A "Filter" button is available. The "Tutorials" section is expanded, showing categories: "TensorFlow tutorials" (Quickstart for beginners, Quickstart for experts), "BEGINNER" (ML basics with Keras, Load and preprocess data), "ADVANCED" (Customization, Distributed training, Images, Text, Audio, Structured data, Generative, Model Understanding, Reinforcement learning). The "Text" category is highlighted with a red box. The "Audio" section is also highlighted with a red box and contains three items: "Simple audio recognition", "Transfer learning for audio recognition", and "Generate music with an RNN". The "Structured data" section contains "Load data" and "Load interface for customization and advanced research." The "Generative" section contains "Neural style transfer", "DeepDream", "DCGAN", "Pix2Pix", "CycleGAN", "Adversarial FGSM", "Intro to Autoencoders", "Variational Autoencoder", and "Lossy data compression". The "Model Understanding" section contains "tf.Estimator" and "Distributed training". A note at the bottom states: "The Advanced section has many instructive notebooks examples, including [Neural machine translation](#), [Transformers](#), and [CycleGAN](#)".

# TensorFlow Tutorials (8/9)



The screenshot shows a TensorFlow tutorial page for a Convolutional Neural Network (CNN) on the TensorFlow Core website. The left sidebar lists various tutorial categories like TensorFlow tutorials, Quickstart for beginners, Beginner ML basics with Keras, Advanced Customization, and more. The main content area is titled "Convolutional Neural Network (CNN)" and includes sections for "Import TensorFlow", "Download and prepare the CIFAR10 dataset", "Verify the data", and "Create the convolutional base". It features code snippets in Google Colab, GitHub links, and download options. Below the code snippets is a grid of 25 CIFAR-10 images with their corresponding labels: frog, truck, truck, deer, automobile, automobile, bird, horse, ship, cat, deer, horse, horse, bird, truck, truck, truck, cat, bird, deer, cat, frog, frog, bird. At the bottom, there's a summary of the model architecture:

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0

<https://www.tensorflow.org/>

# TensorFlow Tutorials (9/9)



TensorFlow > Learn > TensorFlow Core > Tutorials Was this helpful? Up Down

## Simple audio recognition: Recognizing keywords

Run in Google Colab View source on GitHub Download notebook

This tutorial demonstrates how to preprocess audio files in the WAV format and build and train a basic [automatic speech recognition](#) (ASR) model for recognizing ten different words. You will use a portion of the [Speech Commands dataset](#) ([Warden, 2018](#)), which contains short (one-second or less) audio clips of commands, such as "down", "go", "left", "no", "right", "stop", "up" and "yes".

Real-world speech and audio recognition [systems](#) are complex. But, like [image classification with the MNIST dataset](#), this tutorial should give you a basic understanding of the techniques involved.

### Setup

Import necessary modules and dependencies. Note that you'll be using [seaborn](#) for visualization in this tutorial.

```
import os
import pathlib

import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import tensorflow as tf

from tensorflow.keras import layers
from tensorflow.keras import models
from IPython import display

# Set the seed value for experiment reproducibility.
seed = 42
tf.random.set_seed(seed)
np.random.seed(seed)
```

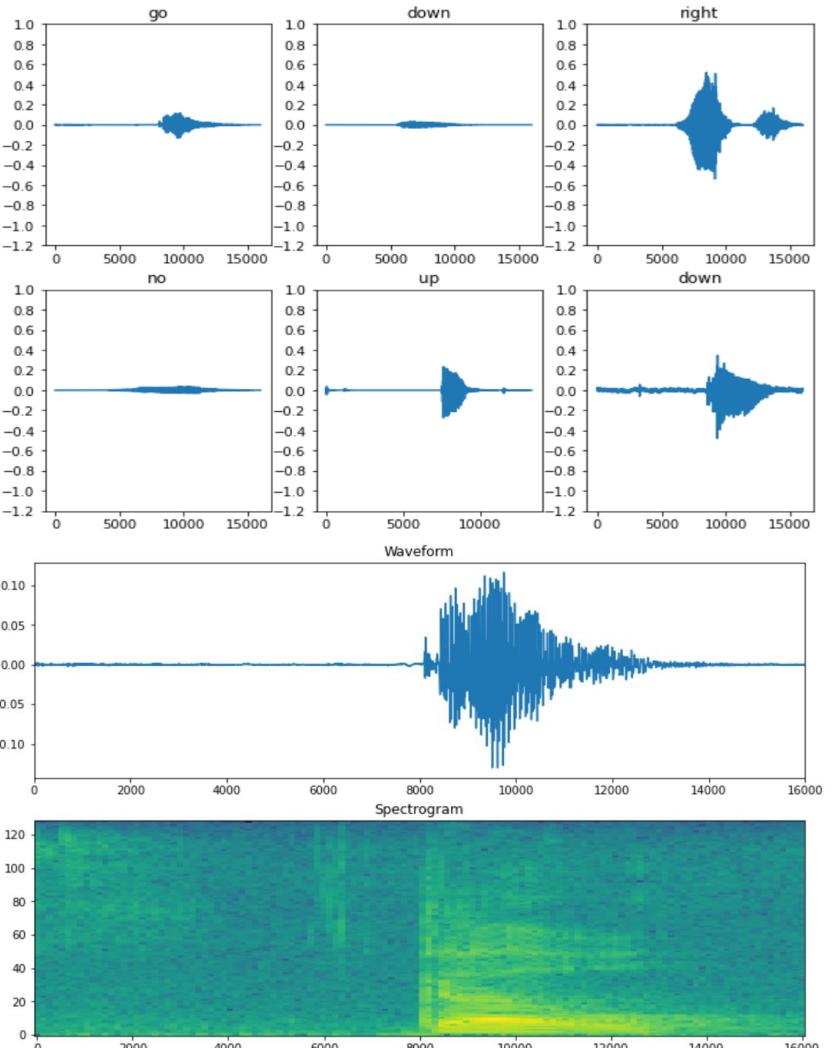
### Import the mini Speech Commands dataset

To save time with data loading, you will be working with a smaller version of the Speech Commands dataset. The [original dataset](#) consists of over 105,000 audio files in the [WAV \(Waveform\) audio file format](#) of people saying 35 different words. This data was collected by Google and released under a CC BY license.

Download and extract the `mini_speech_commands.zip` file containing the smaller Speech Commands datasets with `tf.keras.utils.get_file`:

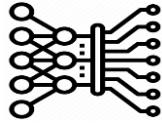
```
DATASET_PATH = 'data/mini_speech_commands'

data_dir = pathlib.Path(DATASET_PATH)
if not data_dir.exists():
    tf.keras.utils.get_file(
        'mini_speech_commands.zip',
        origin="http://storage.googleapis.com/download.tensorflow.org/data/mini_speech_commands.zip",
        extract=True,
        cache_dir='.', cache_subdir='data')
```



<https://www.tensorflow.org/>

# Comparison of libraries



	Keras 	TensorFlow 	PyTorch 
<b>Level of API</b>	high-level API <sup>1</sup>	Both high & low level APIs	Lower-level API <sup>2</sup>
<b>Speed</b>	Slow	High	High
<b>Architecture</b>	Simple, more readable and concise	Not very easy to use	Complex <sup>3</sup>
<b>Debugging</b>	No need to debug	Difficult to debugging	Good debugging capabilities
<b>Dataset Compatibility</b>	Slow & Small	Fast speed & large	Fast speed & large datasets
<b>Popularity Rank</b>	1	2	3
<b>Uniqueness</b>	Multiple back-end support	Object Detection Functionality	Flexibility & Short Training Duration
<b>Created By</b>	Not a library on its own	Created by Google	Created by Facebook <sup>4</sup>
<b>Ease of use</b>	User-friendly	Incomprehensive API	Integrated with Python language
<b>Computational graphs used</b>	Static graphs	Static graphs	Dynamic computation graphs <sup>5</sup>